



---

# MAST6100: MACHINE LEARNING AND DEEP LEARNING

---

Final Project - Report



DECEMBER 1, 2025  
SARAH LEE LWP – WL259

## Table of Contents

Executive Summary .....	3
Introduction .....	4
Background and Motivation .....	4
Aim .....	4
Objectives and Scope.....	4
Sources of Data.....	5
Limitations .....	5
Methodology.....	6
Data Cleaning and Quality Checking .....	6
Exploratory Data Analysis (EDA) .....	6
Variable Selection and Regression .....	6
Classification and Deep Learning .....	7
Analysis and Results .....	9
Conclusion .....	10
Recommendations .....	11
References .....	12
Appendices .....	13
Appendix 1: Installing Packages and Loading Data into R .....	13
Appendix 1.1: Set Working Directory .....	13
Appendix 1.2: Installing Necessary Packages.....	13
Appendix 1.3: Loading Data into R.....	19
Appendix 2: Data Cleaning and Quality Checking .....	20
Appendix 2.1: Loading Necessary Packages .....	20
Appendix 2.2: Quick Data Diagnosis.....	21
Appendix 2.3: Check For Missing Values .....	22
Appendix 2.4: Correcting Letters.....	24
Appendix 2.5: Check For Duplicates .....	24
Appendix 2.6: Checking For Negative Values .....	27
Appendix 2.7: Check The Type Of The Data.....	37
Appendix 2.8: Removing Unneccessary Columns .....	38
Appendix 2.9: Renaming Checked And Cleaned Data Set .....	39
Appendix 3: Exploratory Data Analysis (EDA) .....	40

Appendix 3.1: Loading Necessary Packages .....	40
Appendix 3.2: Basic Statistics and Key Points Of Data.....	40
Appendix 3.3: EDA - Boxplots For Outliers .....	43
Appendix 3.4: EDA - Histogram.....	46
Appendix 3.5: EDA - Pie Charts For All Variables .....	48
Appendix 3.6: EDA - Bar Charts For All Variables .....	53
Appendix 4: Variable Selection And Regression .....	58
Appendix 4.1: Loading Necessary Packages .....	58
Appendix 4.2: Variable Selection.....	59
Appendix 4.3: Prediction .....	63
Appendix 4.4: Variable Selection Interpretation.....	65
Appendix 4.5: Cross-Validation .....	84
Appendix 5: Classification.....	87
Appendix 5.1: Loading Necessary Packages .....	87
Appendix 5.2: Resplit The Data.....	88
Appendix 5.3: Reprocessing And Train Control.....	88
Appendix 5.4: Random Forest.....	88
Appendix 5.5: K-Nearest Neighbors (KNN).....	90
Appendix 5.6: Linear Discriminant Analysis (LDA) .....	91
Appendix 6: Deep Learning .....	92
Appendix 6.1: Loading Necessary Packages .....	92
Appendix 6.1: Resplit The Data.....	93
Appendix 6.2: Reprocessing Data.....	93
Appendix 6.3: Scale Predictors .....	93
Appendix 6.4: Define Neural Network .....	94
Appendix 6.7: Compile And Train Model .....	94
Appendix 6.8: Evaluating Deep Learning Model On Test Data.....	97

## Executive Summary

Hospitals worldwide are increasingly facing pressure from overcrowding, extended waiting times, and rising healthcare costs. Hospital admissions are the pivotal for patients to access timely care and to efficiently use the limited healthcare resources. Inappropriate or delayed admissions can negatively affect patient health and strain hospital capacity which highlights the importance of better understanding the factors that influence admission decisions.

Hospital admissions are shaped by a combination of patients' clinical, demographic, and administrative factors. Gaining insight into how these factors can affect the type of hospital admission can potentially more efficiently support the allocation of resources, improve the delivery of healthcare services, and better inform healthcare policymakers. Motivated by these challenges, this study investigates whether the available patient data can be used to differentiate between the different hospital admission types through data-driven analytical approaches.

The study starts with an exploratory examination of the dataset to assess variable distributions, relationships, and overall data quality. The initial analysis provides an understanding of the underlying structure of the data and supports the subsequent modelling decisions. Then, a range of statistical and machine learning techniques are applied to evaluate the extent to which admission types can be explained or predicted using the given or chosen predictors. Model performance is assessed using standard evaluation metrics and cross-validation to ensure robustness and generalisability of the model.

Instead of using just a single modelling approach, this study implements a more comparative framework by examining both the traditional statistical models and more flexible machine learning and deep learning methods. This approach allows to assess whether an increased model complexity may lead to improvements in performance, or if the data itself has inherent limitations.

Overall, this report provides a logical evaluation of the relationship between available patients' characteristics and hospital admission types. While the findings mainly highlight the challenges of predicting hospital admission types using routinely collected data, they also offer valuable insight into the strengths and limitations of different analytical approaches. These insights can help guide future research, advise data collection strategies, and support healthcare providers and policymakers to improve hospital efficiency and patient care.

# Introduction

Date of report: 1 December 2025

University name: University of Kent

Written by: Sarah Lwp Lee (24022677 – Stage 3 Actuarial Science undergraduate)

This report was prepared by Sarah Lee, in their capacity as a Stage 3 Actuarial Science undergraduate student at the University of Kent with the responsibility of carrying out an analysis on a self-chosen dataset.

Intended recipients(s): Alfred Kume (Lecturer)

## Background and Motivation

Globally, hospitals are currently facing multiple challenges such as overcrowding, extended waiting times and increasing healthcare costs. Amongst these challenges, hospital admissions play a vital role in determining patients' access to immediate medical care. Hence, both avoidable admissions and delayed necessary admissions have emerged as a significant issue in ensuring the appropriate use of hospital services, improving patient outcomes, and maintaining the efficiency of healthcare systems.

Admission decisions can be influenced by multiple factors such as clinical results, demographic and socioeconomic background. Therefore, there is a need for improved understanding of the factors driving hospital admissions.

The motivation for conducting this research is to promote accessible healthcare and more efficiently allocate the limited hospital resources. By identifying and analysing both clinical and non-clinical factors, the findings are expected to inform healthcare providers and policymakers in developing strategies that enhance the quality and efficiency of patient care.

## Aim

The aim of this research is to examine what are the factors that potentially may affect the type of hospital admission (Elective, Urgent or Emergency admissions).

## Objectives and Scope

Analysing patient-related information's association to hospital admission type, including:

- **Age (Years):** Age of the patient at the time of admission, expressed in years.
- **Gender:** Indicates the gender of the patient, either "Male" or "Female."
- **Blood Type:** Patient's blood type
- **Medical Condition:** Primary medical condition or diagnosis associated with the patient, including "Diabetes", "Hypertension", "Asthma", "Obesity", "Arthritis" and "Cancer"

- **Insurance Provider:** Patient's insurance provider, including "Aetna", "Blue Cross", "Cigna", "UnitedHealthcare" and "Medicare"
- **Billing Amount:** The amount of money billed for the patient's healthcare services during their admission
- **Room Number:** The room number of the admitted patient
- **Medication:** The medication prescribed or administered to the patient during their admission, including "Aspirin", "Ibuprofen", "Penicillin", "Paracetamol" and "Lipitor"
- **Test Results:** The results of a medical test conducted during the patient's admission, including "Normal", "Abnormal," or "Inconclusive"

## Sources of Data

In this report, data is primarily sourced from kaggle.com. Additionally, analysis is conducted using:

- Version of R in use: R 4.5.2
- Version of RStudio in use: 2025.09.2+418

To ensure the RMarkdown file runs properly and avoid errors from outdated versions, please use the same version of R and RStudio.

Github link to access all original files: <https://github.com/Sarah-Lwp-Lee/MAST6100--Final-Project.git>

## Limitations

This report is subject to (but not limited to) the following limitations:

- **Methodological Constraints:** Methods used in this report may or may not show or capture all potential relationships
- **Data Quality Issues:**
  - **Data Coverage:** Data only covers admitted patients
  - **Test Results:** Do not know what type of medical test was conducted
- **Time Constraints:** Due to the limited time available, this analysis did not permit a deeper examination and understanding of the data. Hence, the conclusions drawn may not fully reflect all aspects of the dataset.

While this report provides valuable insights, it is important to recognise that some assumptions made during the analysis, such as ignoring the non-existing data, could influence the accuracy of the results and analysis. Therefore, the interpretation of the discussion and analysis should be done with caution due to the inherent limitations of data and the assumptions used.

To address the limitations of this report, please refer to the recommendations section.

# Methodology

## Data Cleaning and Quality Checking

(Please refer to Appendix 2)

Data cleaning and quality checking is used to ensure the data is free of major logical errors and ensure that the data is easy to work with.

Data cleaning and quality checking were done using the following steps:

1. **Quick Data Diagnose (Appendix 2.2):** A preliminary check on the data
2. **Check For Missing Values (Appendix 3.3):** Removing the missing values
3. **Correcting Letters (Appendix 2.4):** For consistency and accuracy in data cleaning and further analysis, all letters are changed to small letters
4. **Check For Duplicates (Appendix 2.5):** To ensure that duplicates will not affect the models, duplicates are removed
5. **Checking For Negative Values (Appendix 2.6):** Some negative values are not logical, so they are removed to prevent it from affecting the models
6. **Check The Type of Data (Appendix 2.7):** Each column's type was changed to its appropriate type so that models will not have errors
7. **Removing Unnecessary Columns (Appendix 2.8):** Some columns were only useful for data cleaning and quality checking, so they are removed

## Exploratory Data Analysis (EDA)

(Please refer to Appendix 3)

Summary of data and multiple types of plots, such as boxplots, histograms, pie charts and bar charts were used to examine the key points and basic patterns in each variable. Correlation between the numeric variables were also used to look for preliminary relationships.

## Variable Selection and Regression

(Please refer to Appendix 4)

From here, admission type is defined as the response variable and all other variables and defined as independent variables (variables that can affect the response variable).

LASSO and Elastic Net was used to select the variables used in the regression model. However, it was proven that the selected variables are not any better than using all the variables. Hence, the full model using all the variables was fitted and used.

The regression model used was multinomial logistic regression as the response variable has 3 categorical outcomes. Dataset was split into 80% training set and 20% testing data. Then, model was fitted using the training set. After that, the fitted model was cross-validated using k-fold cross-validation and the confusion matrix to

determine the accuracy and how much better was the model than just random guessing.

The odds ratio, standard errors and confidence intervals of each independent variable was also considered to determine the degree of effect the variable has on admission types.

## Classification and Deep Learning

(Please refer to Appendix 5 and Appendix 6)

Dataset was split into 80% training set and 20% testing data. Then, instructions were set on how to train and evaluate the models. After that, the training set was used to make the classification models while the test set was used to evaluate the model's accuracy and how much better was the classification model than just random guessing.

The classification models used are:

- **Random Forests**
- **K-Nearest Neighbours (KNN)**
- **Linear Discriminant Analysis (LDA)**

Similarly in deep learning, the dataset was split into 80% training set and 20% testing data. Thereafter, the data was standardised before being used in the neural network.

Neural network has a 30% dropout rate in each of the 2 layers before the final layer to prevent overfitting. Then, the neural network model is compiled and trained while the validation accuracy and validation loss is tracked in the plot below:

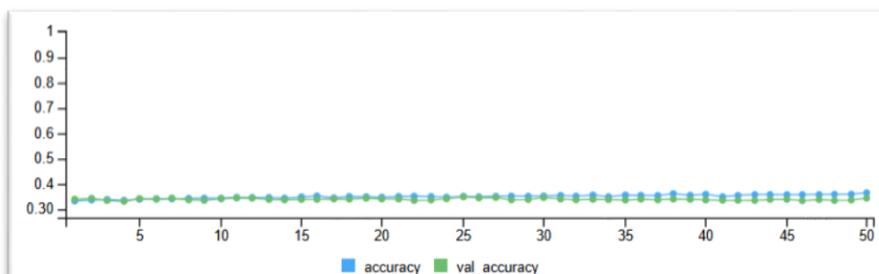


Figure 1: Accuracy Plot

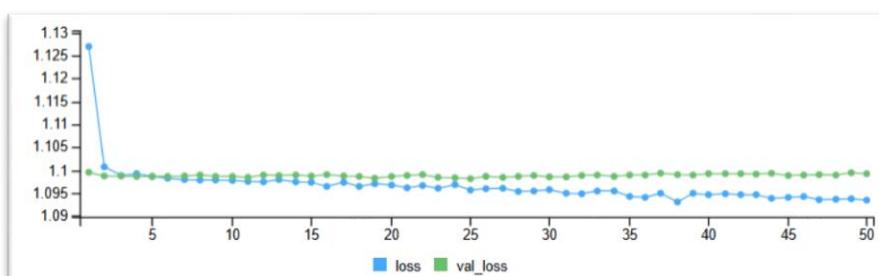


Figure 2: Loss Plot

Then, the neural network is evaluated on the test data. Finally, the overall evaluation is made using the confusion matrix to see the overall accuracy and how much better the deep learning outcomes are than just random guessing.

## Analysis and Results

From the EDA, the dataset shows mostly an even distribution across variables. Numeric variables show no significant correlations, indicating that interaction terms are unnecessary for variable selection later. No outliers were detected. In terms of distribution shape of numeric variables, most are uniformly distributed with only a slight skew. As for the categorical variables, they are all approximately evenly split.

During variable selection, the results indicate that the best-fit model and the full model (using all the independent variables) perform similarly. The residuals indicated a moderate-to-poor fit for the best-fit model. Hence, the full model was used as there were little benefits in using the best-fit model.

After fitting the full regression model, the model shows a low prediction accuracy of around 34%, which was expected given the evenly distributed predictors.

Furthermore, cross-validation confirms the weak predictive power of the model, with 5-fold and 10-fold accuracies being around 33.5-33.8% and Kappa values near zero, which indicates the model performance to be close to random guessing. Exploratory effects only suggest minor patterns with older patients, males, those who are obese, diabetic, asthmatic, certain insurance providers (notably Cigna and Medicare) and some medications show slightly higher likelihoods of emergency and urgent hospital admissions while other predictors have weak or inconsistent effects.

Across the classification models, the performance remains at chance level, indicating no meaningful improvement over the multinomial regression. The random forest model achieves about 33.1% accuracy and has a negative Kappa which gives worse-than-random performance. Similarly, KNN has around 32.9% accuracy and a negative Kappa. LDA only shows a slightly higher accuracy (33.6%) and a marginally positive Kappa (0.0003) effectively no better than random guessing.

Additionally, the neural network model also performs at chance level, with training and validation accuracy at around 33–35%. Referring to Figure 1, the training and validation accuracy remains close which indicates no overfitting of the neural network. Figure 2 support this, as the training loss drops initially but quickly levels off to 1.1 (approximately  $\log(3)$ , the expected loss from a random classification) with the validation loss, shows that the model fails to extract useful patterns from the data. The test accuracy is 33.5% with a Kappa of 0.0092, the model's performance is no better than random guessing. Although the network appears slightly more sensitive to emergency admissions, its low precision leads to frequent misclassifications.

Overall, all models, including the neural network, perform at chance level and fail to reliably predict admission types. Hence, the weak performance indicates that the available predictors lack informative power to meaningfully distinguish between admission types.

## Conclusion

This study assesses whether available patient information could be meaningful in predicting hospital admission types using a range of statistical and machine learning models.

The EDA revealed that the dataset was mostly well-balanced. With numeric variables showing no strong correlations and no outliers, they were approximately uniformly distributed with only slight skewness. Whereas categorical variables were almost evenly split across their respective categories. These characteristics suggested limited inherent structure in the data.

Following that, the fitted multinomial logistic regression model achieved an accuracy of approximately 34%, which was weak but still within expectations. Both the 5-fold and 10-fold cross-validation confirmed the model's weak predictive power. The results indicated that the model's performance was comparable to random guessing. While further exploratory analysis suggested minor tendencies, these effects were minor and inconsistent, offering limited practical predictive value.

When using classification models including random forests, KNN and LDA, all models produced a similar chance-level performance. Random forests and KNN achieved accuracies of approximately 33% with negative Kappa values, indicating worse-than-random predictions. LDA showed a slightly higher accuracy but an effectively zero Kappa, showing no meaningful improvement over other approaches or the baseline regression model.

Finally, a neural network model was built to determine whether more complex, non-linear relationships could be learned from the data. The neural network also performed at the chance level and showed that predictions was no better than random guessing even though there were no indications of overfitting. Although the model did show a slightly higher sensitivity to emergency admissions, its low precision resulted in frequent misclassifications.

In summary, all the models evaluated in this study failed to predict hospital admission types beyond the chance level. This consistent pattern strongly suggests that the primary limitation lies not in model choice or complexity, but that the available variables lack informative power to capture the signal to distinguish between elective, urgent and emergency hospital admissions. Further research would be required with other clinically relevant features, such as detailed medical history or severity indicators, to improve the models' predictive performance.

## Recommendations

1. **Improve Data Collection and Integration Practices:** Healthcare systems should invest in better integration of health records and clinical databases. Improved collection of data (e.g. knowing what type of medical test was conducted) can form a more comprehensive patient profile, allowing more effective analytical and predictive modelling.
2. **Enhance Data Quality and Clinical Detail:** Future analyses should include other clinical variables that better reflect the patients' status. These factors may provide stronger predictive signals than just demographic or administrative variables alone.
3. **Expand Socioeconomic and Access-to-Care Variables:** While basic demographic variables were included, more detailed socioeconomic indicators (e.g. income, employment status or access to primary care) may help explain admission patterns, particularly avoidable or delayed admissions.
4. **Reconsider Outcome Definition and Class Structure:** The current admission type categories may be too broad. Future studies could explore alternative groupings (e.g. emergency vs non-emergency) or hierarchical classification approaches to better capture distinctions in admission decisions.

By implementing these recommendations, future reports and analysis could be more reliable, accurate and actionable.

## References

- Kaggle. (n.d.). *Healthcare dataset* [Data set]. Kaggle. Retrieved December 17, 2025, from <https://www.kaggle.com/datasets/prasad22/healthcare-dataset>
- *Finding duplicate values in a data frame in R: A guide using base R and dplyr.* (2023, July 16). R-bloggers. <https://www.r-bloggers.com/2023/07/finding-duplicate-values-in-a-data-frame-in-r-a-guide-using-base-r-and-dplyr/>
- *Identify and remove duplicate data in R.* (2018, October 19). Datanovia. <https://www.datanovia.com/en/lessons/identify-and-remove-duplicate-data-in-r/>
- (n.d.). The Comprehensive R Archive Network. <https://cran.r-project.org/web/packages/e1071/e1071.pdf>
- (n.d.). The Comprehensive R Archive Network. <https://cran.r-project.org/web/packages/tidyverse/tidyverse.pdf>
- (n.d.). The Comprehensive R Archive Network. <https://cran.r-project.org/web/packages/dlookr/dlookr.pdf>
- (n.d.). The Comprehensive R Archive Network. <https://cran.r-project.org/web/packages/dplyr/dplyr.pdf>
- (n.d.). The Comprehensive R Archive Network. <https://cran.r-project.org/web/packages/ggplot2/ggplot2.pdf>
- (n.d.). The Comprehensive R Archive Network. <https://cran.r-project.org/web/packages/tidyr/tidyr.pdf>
- *A Short Introduction to the caret Package.* (n.d.). <https://cran.r-project.org/web/packages/caret/vignettes/caret.html>
- Liaw, A. & Matthew Wiener. (2024). *Breiman and Cutlers random forests for classification and regression.* CRAN. <https://cran.r-project.org/web/packages/randomForest/randomForest.pdf>
- *Package ‘kernlab.’* (2025). <https://cran.r-project.org/web/packages/kernlab/kernlab.pdf>
- *Getting Started with Keras.* (n.d.). [https://cran.r-project.org/web/packages/keras3/vignettes/getting\\_started.html](https://cran.r-project.org/web/packages/keras3/vignettes/getting_started.html)
- Allaire, J., Kalinowski, T., Falbel, D., Eddelbuettel, D., Yuan Tang, Golding, N., & Google Inc. (2025). *Package ‘tensorflow.’* CRAN. <https://cran.r-project.org/web/packages/tensorflow/tensorflow.pdf>
- Ripley, B., Venables, W., & R-project.org. (2025). *Functions for classification* [Functions for Classification]. <https://cran.r-project.org/web/packages/class/class.pdf>
- Silge, J., & Robinson, D. (2023, January 7). *Introduction to tidytext.* The Comprehensive R Archive Network. <https://cran.r-project.org/web/packages/tidytext/vignettes/tidytext.html>
- McQuire, P., & Kume, A. (2023). *R programming for actuarial science.* John Wiley & Sons.

# Appendices

## Appendix 1: Installing Packages and Loading Data into R

### Appendix 1.1: Set Working Directory

Setting the working directory enables R where to retrieve the data file from, so that we can load the data set for analysis.

Set working directory of where the data file is located by: ‘Session’ → ‘Set Working Directory’ → ‘Choose Directory’ → choose the file of where the data file is located OR use the code below by inserting the file’s location and removing the # to run the code.

```
# setwd("insert_file_location_here")
```

### Appendix 1.2: Installing Necessary Packages

Before we start, we need to install the packages used in each section to use the functions we need for analysis. Later on in each section, we will load the packages used for the section.

Note: Some packages are used in different sections as well, but a package only needs to be installed once.

```
# Data Cleaning and Quality Checking
install.packages("tidyverse", dependencies = TRUE, repos = "http://cran.rstudio.com")

## Installing package into 'C:/Users/Acer User/AppData/Local/R/win-library/4.5'
## (as 'lib' is unspecified)

## package 'tidyverse' successfully unpacked and MD5 sums checked
##
## The downloaded binary packages are in
## C:\Users\Acer User\AppData\Local\Temp\RtmpCGvgyq\downloaded_packages

install.packages("dlookr", dependencies = TRUE, repos = "http://cran.rstudio.com")

## Installing package into 'C:/Users/Acer User/AppData/Local/R/win-library/4.5'
## (as 'lib' is unspecified)

## package 'dlookr' successfully unpacked and MD5 sums checked
##
## The downloaded binary packages are in
## C:\Users\Acer User\AppData\Local\Temp\RtmpCGvgyq\downloaded_packages

# Exploratory Data Analysis (EDA)
install.packages("dplyr", dependencies = TRUE, repos = "http://cran.rstudio.com")
```

```

## Installing package into 'C:/Users/Acer User/AppData/Local/R/win-library
/4.5'
## (as 'lib' is unspecified)

## package 'dplyr' successfully unpacked and MD5 sums checked

## Warning: cannot remove prior installation of package 'dplyr'

## Warning in file.copy(savedcopy, lib, recursive = TRUE): problem copying
## C:\Users\Acer
## User\AppData\Local\R\win-library\4.5\00LOCK\dplyr\libs\x64\dplyr.dll to
## C:\Users\Acer User\AppData\Local\R\win-library\4.5\dplyr\libs\x64\dplyr
.dll:
## Permission denied

## Warning: restored 'dplyr'

##
## The downloaded binary packages are in
## C:\Users\Acer User\AppData\Local\Temp\RtmpCGvgyq\downloaded_packages

install.packages("ggplot2", dependencies = TRUE, repos = "http://cran.rstudio.com")

## Installing package into 'C:/Users/Acer User/AppData/Local/R/win-library
/4.5'
## (as 'lib' is unspecified)

## package 'ggplot2' successfully unpacked and MD5 sums checked
##
## The downloaded binary packages are in
## C:\Users\Acer User\AppData\Local\Temp\RtmpCGvgyq\downloaded_packages

install.packages("tidyverse", dependencies = TRUE, repos = "http://cran.rstudio.com")

## Installing package into 'C:/Users/Acer User/AppData/Local/R/win-library
/4.5'
## (as 'lib' is unspecified)

## package 'tidyverse' successfully unpacked and MD5 sums checked

## Warning: cannot remove prior installation of package 'tidyverse'

## Warning in file.copy(savedcopy, lib, recursive = TRUE): problem copying
## C:\Users\Acer
## User\AppData\Local\R\win-library\4.5\00LOCK\tidyverse\libs\x64\tidyverse.dll to
## C:\Users\Acer User\AppData\Local\R\win-library\4.5\tidyverse\libs\x64\tidyverse
.dll:
## Permission denied

## Warning: restored 'tidyverse'

##
## The downloaded binary packages are in
## C:\Users\Acer User\AppData\Local\Temp\RtmpCGvgyq\downloaded_packages

```

```

install.packages("scales", dependencies = TRUE, repos = "http://cran.rstudio.com")

## Installing package into 'C:/Users/Acer User/AppData/Local/R/win-library/4.5'
## (as 'lib' is unspecified)

## package 'scales' successfully unpacked and MD5 sums checked
##
## The downloaded binary packages are in
## C:\Users\Acer User\AppData\Local\Temp\RtmpCGvgyq\downloaded_packages

install.packages("janitor", dependencies = TRUE, repos = "http://cran.rstudio.com")

## Installing package into 'C:/Users/Acer User/AppData/Local/R/win-library/4.5'
## (as 'lib' is unspecified)

## package 'janitor' successfully unpacked and MD5 sums checked
##
## The downloaded binary packages are in
## C:\Users\Acer User\AppData\Local\Temp\RtmpCGvgyq\downloaded_packages

# Variable Selection and Regression
install.packages("MASS", dependencies = TRUE, repos = "http://cran.rstudio.com")

## Installing package into 'C:/Users/Acer User/AppData/Local/R/win-library/4.5'
## (as 'lib' is unspecified)

## package 'MASS' successfully unpacked and MD5 sums checked

## Warning: cannot remove prior installation of package 'MASS'

## Warning in file.copy(savedcopy, lib, recursive = TRUE): problem copying
## C:\Users\Acer
## User\AppData\Local\R\win-library\4.5\00LOCK\MASS\libs\x64\MASS.dll to
## C:\Users\Acer User\AppData\Local\R\win-library\4.5\MASS\libs\x64\MASS.dll:
## Permission denied

## Warning: restored 'MASS'

##
## The downloaded binary packages are in
## C:\Users\Acer User\AppData\Local\Temp\RtmpCGvgyq\downloaded_packages

install.packages("glmnet", dependencies = TRUE, repos = "http://cran.rstudio.com")

## Installing package into 'C:/Users/Acer User/AppData/Local/R/win-library/4.5'
## (as 'lib' is unspecified)

## package 'glmnet' successfully unpacked and MD5 sums checked

```

```

## Warning: cannot remove prior installation of package 'glmnet'

## Warning in file.copy(savedcopy, lib, recursive = TRUE): problem copying
## C:\Users\Acer
## User\AppData\Local\R\win-library\4.5\00LOCK\glmnet\libs\x64\glmnet.dll
## C:\Users\Acer User\AppData\Local\R\win-library\4.5\glmnet\libs\x64\glmnet.dll:
## Permission denied

## Warning: restored 'glmnet'

##
## The downloaded binary packages are in
## C:\Users\Acer User\AppData\Local\Temp\RtmpCGvgyq\downloaded_packages

install.packages("Matrix", dependencies = TRUE, repos = "http://cran.rstudio.com")

## Installing package into 'C:/Users/Acer User/AppData/Local/R/win-library/4.5'
## (as 'lib' is unspecified)

## package 'Matrix' successfully unpacked and MD5 sums checked

## Warning: cannot remove prior installation of package 'Matrix'

## Warning in file.copy(savedcopy, lib, recursive = TRUE): problem copying
## C:\Users\Acer
## User\AppData\Local\R\win-library\4.5\00LOCK\Matrix\libs\x64\Matrix.dll
## C:\Users\Acer User\AppData\Local\R\win-library\4.5\Matrix\libs\x64\Matrix.dll:
## Permission denied

## Warning: restored 'Matrix'

##
## The downloaded binary packages are in
## C:\Users\Acer User\AppData\Local\Temp\RtmpCGvgyq\downloaded_packages

install.packages("nnet", dependencies = TRUE, repos = "http://cran.rstudio.com")

## Installing package into 'C:/Users/Acer User/AppData/Local/R/win-library/4.5'
## (as 'lib' is unspecified)

## package 'nnet' successfully unpacked and MD5 sums checked

## Warning: cannot remove prior installation of package 'nnet'

## Warning in file.copy(savedcopy, lib, recursive = TRUE): problem copying
## C:\Users\Acer
## User\AppData\Local\R\win-library\4.5\00LOCK\nnet\libs\x64\nnet.dll to
## C:\Users\Acer User\AppData\Local\R\win-library\4.5\nnet\libs\x64\nnet.d

```

```

11:
## Permission denied

## Warning: restored 'nnet'

##
## The downloaded binary packages are in
## C:\Users\Acer User\AppData\Local\Temp\RtmpCGvgyq\downloaded_packages

install.packages("caret", dependencies = TRUE, repos = "http://cran.rstudio.com")

## Installing package into 'C:/Users/Acer User/AppData/Local/R/win-library/4.5'
## (as 'lib' is unspecified)

## package 'caret' successfully unpacked and MD5 sums checked

## Warning: cannot remove prior installation of package 'caret'

## Warning in file.copy(savedcopy, lib, recursive = TRUE): problem copying
## C:\Users\Acer
## User\AppData\Local\R\win-library\4.5\00LOCK\caret\libs\x64\caret.dll to
## C:\Users\Acer User\AppData\Local\R\win-library\4.5\caret\libs\x64\caret.dll:
## Permission denied

## Warning: restored 'caret'

##
## The downloaded binary packages are in
## C:\Users\Acer User\AppData\Local\Temp\RtmpCGvgyq\downloaded_packages

install.packages("hnp", dependencies = TRUE, repos = "http://cran.rstudio.com")

## Installing package into 'C:/Users/Acer User/AppData/Local/R/win-library/4.5'
## (as 'lib' is unspecified)

## Warning: dependency 'glmmADMB' is not available

## package 'hnp' successfully unpacked and MD5 sums checked
##
## The downloaded binary packages are in
## C:\Users\Acer User\AppData\Local\Temp\RtmpCGvgyq\downloaded_packages

# Classification
install.packages("e1071", dependencies = TRUE, repos = "http://cran.rstudio.com")

## Installing package into 'C:/Users/Acer User/AppData/Local/R/win-library/4.5'
## (as 'lib' is unspecified)

## package 'e1071' successfully unpacked and MD5 sums checked

```

```

## Warning: cannot remove prior installation of package 'e1071'

## Warning in file.copy(savedcopy, lib, recursive = TRUE): problem copying
## C:\Users\Acer
## User\AppData\Local\R\win-library\4.5\00LOCK\e1071\libs\x64\e1071.dll to
## C:\Users\Acer User\AppData\Local\R\win-library\4.5\e1071\libs\x64\e1071
.dll:
## Permission denied

## Warning: restored 'e1071'

## 
## The downloaded binary packages are in
## C:\Users\Acer User\AppData\Local\Temp\RtmpCGvgyq\downloaded_packages

install.packages("randomForest", dependencies = TRUE, repos = "http://cran.rstudio.com")

## Installing package into 'C:/Users/Acer User/AppData/Local/R/win-library
/4.5'
## (as 'lib' is unspecified)

## Warning: package 'randomForest' is not available for this version of R
##
## A version of this package for your version of R might be available elsewhere,
## see the ideas at
## https://cran.r-project.org/doc/manuals/r-patched/R-admin.html#Installing-packages

install.packages("class", dependencies = TRUE, repos = "http://cran.rstudio.com")

## Installing package into 'C:/Users/Acer User/AppData/Local/R/win-library
/4.5'
## (as 'lib' is unspecified)

## package 'class' successfully unpacked and MD5 sums checked

## Warning: cannot remove prior installation of package 'class'

## Warning in file.copy(savedcopy, lib, recursive = TRUE): problem copying
## C:\Users\Acer
## User\AppData\Local\R\win-library\4.5\00LOCK\class\libs\x64\class.dll to
## C:\Users\Acer User\AppData\Local\R\win-library\4.5\class\libs\x64\class
.dll:
## Permission denied

## Warning: restored 'class'

## 
## The downloaded binary packages are in
## C:\Users\Acer User\AppData\Local\Temp\RtmpCGvgyq\downloaded_packages

install.packages("kernlab", dependencies = TRUE, repos = "http://cran.rstudio.com")

```

```

## Installing package into 'C:/Users/Acer User/AppData/Local/R/win-library
/4.5'
## (as 'lib' is unspecified)

## package 'kernlab' successfully unpacked and MD5 sums checked

## Warning: cannot remove prior installation of package 'kernlab'

## Warning in file.copy(savedcopy, lib, recursive = TRUE): problem copying
## C:\Users\Acer
## User\AppData\Local\R\win-library\4.5\00LOCK\kernlab\libs\x64\kernlab.dl
1 to
## C:\Users\Acer
## User\AppData\Local\R\win-library\4.5\kernlab\libs\x64\kernlab.dll: Perm
ission
## denied

## Warning: restored 'kernlab'

##
## The downloaded binary packages are in
## C:\Users\Acer User\AppData\Local\Temp\RtmpCGvgyq\downloaded_packages

# Deep Learning
install.packages("keras3", dependencies = TRUE, repos = "http://cran.rstud
io.com")

## Installing package into 'C:/Users/Acer User/AppData/Local/R/win-library
/4.5'
## (as 'lib' is unspecified)

## package 'keras3' successfully unpacked and MD5 sums checked
##
## The downloaded binary packages are in
## C:\Users\Acer User\AppData\Local\Temp\RtmpCGvgyq\downloaded_packages

install.packages("tensorflow", dependencies = TRUE, repos = "http://cran.r
studio.com")

## Installing package into 'C:/Users/Acer User/AppData/Local/R/win-library
/4.5'
## (as 'lib' is unspecified)

## package 'tensorflow' successfully unpacked and MD5 sums checked
##
## The downloaded binary packages are in
## C:\Users\Acer User\AppData\Local\Temp\RtmpCGvgyq\downloaded_packages

```

## Appendix 1.3: Loading Data into R

Load the data file into RStudio using the code below and look at the beginning of the data to see some features of the data. (E.g.: what are the variable names in the file, the types of data)

```
# To Load the data into R for a csv file first remove #, then insert file
name:
```

```

# raw_data <- read.csv("insert_name_of_your_data.csv")
raw_data <- read.csv("healthcare_dataset.csv") # Name the Loaded data as "raw data" as the data has not been checked

head(raw_data) # To show beginning of data

##          Name Age Gender Blood.Type Medical.Condition Date.of.Admission
## 1 Bobby JacksOn 30   Male      B-        Cancer 2024-01-31
## 2 LesLie TErRy 62   Male      A+        Obesity 2019-08-20
## 3 DaNnY sMith 76 Female     A-        Obesity 2022-09-22
## 4 andrEw waTTs 28 Female     O+        Diabetes 2020-11-18
## 5 adrIENNE bEll 43 Female     AB+       Cancer 2022-09-19
## 6 EMILY JOHNSOn 36   Male      A+        Asthma 2023-12-20
##          Doctor                               Hospital Insurance.Provider Billings.Amount
## 1 Matthew Smith           Sons and Miller      Blue Cross 18856.28
## 2 Samantha Davies         Kim Inc            Medicare 33643.33
## 3 Tiffany Mitchell        Cook PLC           Aetna 27955.10
## 4 Kevin Wells Hernandez Rogers and Vang,    Medicare 37909.78
## 5 Kathleen Hanna          White-White        Aetna 14238.32
## 6 Taylor Newton           Nunez-Humphrey    UnitedHealthcare 48145.11
##          Room.Number Admission.Type Discharge.Date Medication Test.Results
## 1          328      Urgent 2024-02-02 Paracetamol      Normal
## 2          265      Emergency 2019-08-26 Ibuprofen Inconclusive
## 3          205      Emergency 2022-10-07 Aspirin      Normal
## 4          450      Elective 2020-12-18 Ibuprofen     Abnormal
## 5          458      Urgent 2022-10-09 Penicillin     Abnormal
## 6          389      Urgent 2023-12-24 Ibuprofen      Normal

```

## Appendix 2: Data Cleaning and Quality Checking

### Appendix 2.1: Loading Necessary Packages

```

library(tidyverse) # For data manipulation, visualization and analysis

## — Attaching core tidyverse packages —————— tidyverse
2.0.0 —
## ✓ dplyr     1.1.4    ✓ readr     2.1.5
## ✓ forcats   1.0.1    ✓ stringr   1.5.2
## ✓ ggplot2   4.0.1    ✓ tibble    3.3.0
## ✓ lubridate 1.9.4    ✓ tidyverse  1.3.1

```

```

## ✓ purrr      1.1.0
## — Conflicts _____ tidyverse_conflicts() —
## X dplyr::filter() masks stats::filter()
## X dplyr::lag()    masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force
## all conflicts to become errors

library(dlookr)    # This allows us to create a table with the number of m
issing values and gives us other insights

## Registered S3 methods overwritten by 'dlookr':
##   method      from
##   plot.transform scales
##   print.transform scales
##
## Attaching package: 'dlookr'
##
## The following object is masked from 'package:tidyr':
## 
##   extract
##
## The following object is masked from 'package:base':
## 
##   transform

```

## Appendix 2.2: Quick Data Diagnosis

In this section, we will do a preliminary check on the data to see what might need to be done in data cleaning and quality checking.

```

str(raw_data) # Checking the structure of data frame

## 'data.frame': 55500 obs. of 15 variables:
## $ Name          : chr "Bobby JacksOn" "LesLie TErRy" "DaNnY sMitH"
## "andrEw waTtS" ...
## $ Age           : int 30 62 76 28 43 36 21 20 82 58 ...
## $ Gender        : chr "Male" "Male" "Female" "Female" ...
## $ Blood.Type    : chr "B-" "A+" "A-" "O+" ...
## $ Medical.Condition: chr "Cancer" "Obesity" "Obesity" "Diabetes" ...
## $ Date.of.Admission: chr "2024-01-31" "2019-08-20" "2022-09-22" "202
0-11-18" ...
## $ Doctor         : chr "Matthew Smith" "Samantha Davies" "Tiffany
Mitchell" "Kevin Wells" ...
## $ Hospital       : chr "Sons and Miller" "Kim Inc" "Cook PLC" "Her
nandez Rogers and Vang," ...
## $ Insurance.Provider: chr "Blue Cross" "Medicare" "Aetna" "Medicare"
...
## $ Billing.Amount  : num 18856 33643 27955 37910 14238 ...
## $ Room.Number     : int 328 265 205 450 458 389 389 277 316 249 ...
## $ Admission.Type  : chr "Urgent" "Emergency" "Emergency" "Elective"
...
## $ Discharge.Date  : chr "2024-02-02" "2019-08-26" "2022-10-07" "202
0-12-18" ...
## $ Medication      : chr "Paracetamol" "Ibuprofen" "Aspirin" "Ibupro

```

```

fen" ...
## $ Test.Results      : chr  "Normal" "Inconclusive" "Normal" "Abnormal"
...
dim(raw_data) # Checking the dimensions of the data
## [1] 55500     15

diagnosis <- diagnose(raw_data) # This creates a table of variable names,
# the type of data, number of missing values and unique values, along with t
he rate of unique values
diagnosis # This allows us to see the whole diagnosis.

## # A tibble: 15 × 6
##   variables      types missing_count missing_percent unique_count un
ique_rate
##   <chr>        <chr>        <int>             <dbl>        <int>
## 1 Name          char...         0                 0       49992
0.901
## 2 Age           inte...         0                 0        77
0.00139
## 3 Gender         char...         0                 0        2
0.0000360
## 4 Blood.Type    char...         0                 0        8
0.000144
## 5 Medical.Conditi... char...         0                 0        6
0.000108
## 6 Date.of.Admissi... char...         0                 0      1827
0.0329
## 7 Doctor         char...         0                 0      40341
0.727
## 8 Hospital        char...         0                 0      39876
0.718
## 9 Insurance.Provi... char...         0                 0        5
0.0000901
## 10 Billing.Amount  nume...         0                 0      50000
0.901
## 11 Room.Number    inte...         0                 0        400
0.00721
## 12 Admission.Type char...         0                 0        3
0.0000541
## 13 Discharge.Date char...         0                 0      1856
0.0334
## 14 Medication      char...         0                 0        5
0.0000901
## 15 Test.Results    char...         0                 0        3
0.0000541

```

## Appendix 2.3: Check For Missing Values

There may be certain values in the data that are not recognised as missing values, so we will first replace all these possible values. Then, we can see the true missing data count to determine if the data needs to be cleaned for missing values.

```

raw_data[raw_data == "NA"] <- NA
# If there are any values that says NA but recognises it as text, it might
# cause the missing count to be 0, so this fixes this issue
raw_data[raw_data == "[ ]"] <- NA
# If there are any values that says [] but recognises it as text, it might
# cause the missing count to be 0, so this fixes this issue
raw_data[raw_data == ""] <- NA
# If there are any values that says empty cells but recognises it as text,
# it might cause the missing count to be 0, so this fixes this issue

# Now we can see the true missing data count in diagnose again
diagnosis <- diagnose(raw_data)
diagnosis

## # A tibble: 15 × 6
##   variables      types missing_count missing_percent unique_count un
##   <chr>          <chr>     <int>           <dbl>        <int>
## 1 Name            char...       0             0         49992
## 2 Age             inte...       0             0           77
## 3 Gender          char...       0             0             2
## 4 Blood.Type      char...       0             0             8
## 5 Medical.Conditi... char...       0             0             6
## 6 Date.of.Admissi... char...       0             0           1827
## 7 Doctor          char...       0             0           40341
## 8 Hospital         char...       0             0           39876
## 9 Insurance.Provi... char...       0             0             5
## 10 Billing.Amount   nume...       0             0           50000
## 11 Room.Number     inte...       0             0             400
## 12 Admission.Type   char...       0             0             3
## 13 Discharge.Date   char...       0             0           1856
## 14 Medication       char...       0             0             5
## 15 Test.Results     char...       0             0             3
## # ... with 15 more variables:
## #   .linenos. = c(1L, 2L, 3L, 4L, 5L, 6L, 7L, 8L, 9L, 10L, 11L, 12L, 13L, 14L, 15L, 16L),
## #   .type. = c("tbl_df", "tbl_structured_text", "tbl_header", "tbl_body", "tbl_header", "tbl_body", "tbl_header", "tbl_body", "tbl_header", "tbl_body", "tbl_header", "tbl_body", "tbl_header", "tbl_body", "tbl_header", "tbl_body"),
## #   .format. = c("text", "text", "text"),
## #   .missing. = c("no", "no", "no"),
## #   .source. = c("text", "text", "text"),
## #   .order. = c(1L, 2L, 3L, 4L, 5L, 6L, 7L, 8L, 9L, 10L, 11L, 12L, 13L, 14L, 15L, 16L),
## #   .width. = c(10, 10, 10, 10, 10, 10, 10, 10, 10, 10, 10, 10, 10, 10, 10, 10),
## #   .label. = c("variables", "types", "missing_count", "missing_percent", "unique_count", ".linenos.", ".type.", ".format.", ".missing.", ".source.", ".order.", ".width.", ".label."))

# We can see that there are no missing observations in any of the data. To
# confirm this, we can use another line of code as below:
sum(is.na(raw_data)) # Shows total number of missing observations, which is 0 and confirms our diagnosis above

```

```
## [1] 0
```

## Appendix 2.4: Correcting Letters

We can see from the beginning of data that the patient names do not have consistent letters. So, to make data cleaning more accurate, we will make all the letters in the name into small letters.

```
# Changing patient names to all small letters
raw_data$Name <- tolower(raw_data$Name)
```

## Appendix 2.5: Check For Duplicates

Check if all data in the dataset is unique and there are no repeats.

If all data taken is unique, then the patients' names should be all unique and match the number of rows. So, we will check for the names of all the patients and ensure there are no repeats.

```
# Step 1: Check if all patient names are unique
# If all patient names are unique, then the number of unique rows should match with the number of rows of the data
if(n_distinct(raw_data$Name) == nrow(raw_data)) {
  print(paste("No duplicates in patient names"))
} else {
  num_duplicates <- nrow(raw_data)-n_distinct(raw_data$Name)
  print(paste("Number of potential duplicates in patient names:", num_duplicates))
}

## [1] "Number of potential duplicates in patient names: 15265"

# Output shows that there is a duplicate patient names, so we are going to identify it using the duplicated function with a condition to show the first duplicate line. Then, we match the admission data and discharge date of the patients with the same name to identify the difference between these duplicated names.

duplicated_rows1 <- raw_data %>%
  filter(
    duplicated(across(c(Name))) |
    duplicated(across(c(Name,)), fromLast = TRUE)
  ) %>%
  arrange(Name)

print(head(duplicated_rows1))

##           Name Age Gender Blood.Type Medical.Condition Date.of.Admission
## 1   aaron archer  47 Female      B-        Cancer 2021-01
## 2   aaron archer  49 Female      B-        Cancer 2021-01
## 3   aaron baker  73   Male      B+        Cancer 2019-06
## 4   aaron baker  73   Male      B+        Cancer 2019-06
## 5   aaron baker  73   Male      B+        Cancer 2019-06
## 6   aaron baker  73   Male      B+        Cancer 2019-06
```

```

## 4 aaron baker 84 Female A+ Asthma 2022-06
-04
## 5 aaron bradshaw 25 Female 0+ Arthritis 2019-11
-30
## 6 aaron bradshaw 78 Male B- Arthritis 2019-11
-15
## Doctor Hospital Insurance.Provider
## 1 Cynthia Villanueva Montes Case and Mendez, Medicare
## 2 Cynthia Villanueva Montes Case and Mendez, Medicare
## 3 Tracy Torres Wise and Todd, Parker Medicare
## 4 Brittany Smith Carter, Abbott and Fuentes Medicare
## 5 Mackenzie Phillips Olson LLC Aetna
## 6 Sheila Smith Group Martinez UnitedHealthcare
## Billing.Amount Room.Number Admission.Type Discharge.Date Medication
## 1 10602.077 108 Urgent 2021-01-17 Paracetamol
## 2 10602.077 108 Urgent 2021-01-17 Paracetamol
## 3 10135.885 234 Elective 2019-07-10 Aspirin
## 4 6826.677 496 Emergency 2022-06-10 Lipitor
## 5 16342.364 255 Emergency 2019-12-20 Lipitor
## 6 50132.996 393 Emergency 2019-11-17 Penicillin
## Test.Results
## 1 Inconclusive
## 2 Inconclusive
## 3 Inconclusive
## 4 Normal
## 5 Inconclusive
## 6 Abnormal

# Here, we can see that most of the duplicated names, but it is hard to tell which to remove
# So, we will filter it again to look for the patients that have the exact same details for all their columns. This is to enable us to know which of them are definitely duplicates that need to be removed
exact_duplicates <- raw_data %>%
  group_by(across(everything())) %>%
  filter(n() > 1) %>%
  ungroup() %>%
  arrange(Name)

print(head(exact_duplicates))

## # A tibble: 6 × 15
##   Name          Age Gender Blood.Type Medical.Condition Date.of.Admission
##   <chr>        <int> <chr>  <chr>    <chr>           <chr>
## 1 abigail y...     41 Female 0+    Hypertension 2022-12-15
## 2 abigail y...     41 Female 0+    Hypertension 2022-12-15
## 3 adam thomas      75 Male   0+    Hypertension 2022-01-02
## 4 adam thomas      75 Male   0+    Hypertension 2022-01-02

```

```

## 5 alex black      51 Male   0+       Diabetes      2022-03-27
# Frank...
## 6 alex black      51 Male   0+       Diabetes      2022-03-27
# Frank...
## # i 8 more variables: Hospital <chr>, Insurance.Provider <chr>,
## #   Billing.Amount <dbl>, Room.Number <int>, Admission.Type <chr>,
## #   Discharge.Date <chr>, Medication <chr>, Test.Results <chr>

# Here, we can see that there are 1068/2=534 pairs of duplicated names with the exact same details, so we will remove one of each

raw_data <- raw_data %>%
  distinct()
# To check if this was successful, we will run the exact_duplicateds again.
# If it prints to have no rows, we will know that this is successful

exact_duplicates <- raw_data %>%
  group_by(across(everything())) %>%
  filter(n() > 1) %>%
  ungroup() %>%
  arrange(Name)

print(head(exact_duplicates))

## # A tibble: 0 × 15
## # i 15 variables: Name <chr>, Age <int>, Gender <chr>, Blood.Type <chr>
# ,
## #   Medical.Condition <chr>, Date.of.Admission <chr>, Doctor <chr>,
## #   Hospital <chr>, Insurance.Provider <chr>, Billing.Amount <dbl>,
## #   Room.Number <int>, Admission.Type <chr>, Discharge.Date <chr>,
## #   Medication <chr>, Test.Results <chr>

# There are no rows printed, so the removal was successful

# Now, we will check if the patients with the same name have any different
# details on each column
other_cols <- setdiff(names(raw_data), "Name") # Get all columns except Name

name_duplicates <- raw_data %>%
  group_by(Name) %>%
  # Check if any other column varies within the same Name
  filter(if_any(all_of(other_cols), ~ n_distinct(.) > 1)) %>%
  ungroup() %>%
  arrange(Name)

print(head(name_duplicates))

## # A tibble: 6 × 15
##   Name          Age Gender Blood.Type Medical.Condition Date.of.Admission
##   <chr>        <int> <chr>  <chr>    <chr>           <chr>
## 1 aaron arch...     47 Female B-
#   Cynth...

```

```

## 2 aaron arch...    49 Female B-      Cancer        2021-01-10
  Cynth...
## 3 aaron baker     73 Male   B+      Cancer        2019-06-18
  Tracy...
## 4 aaron baker     84 Female A+     Asthma       2022-06-04
  Britt...
## 5 aaron brad...    25 Female O+     Arthritis    2019-11-30
  Macke...
## 6 aaron brad...    78 Male   B-      Arthritis    2019-11-15
  Sheil...
## # i 8 more variables: Hospital <chr>, Insurance.Provider <chr>,
## #   Billing.Amount <dbl>, Room.Number <int>, Admission.Type <chr>,
## #   Discharge.Date <chr>, Medication <chr>, Test.Results <chr>

# From what we can observe, most of the patients with the same name have a
# almost the same details except for their age. But since we cannot tell what
# is the real age of the patient, we will assume the first entry of the per-
# son in the data is correct as there are no missing data for each column.
raw_data <- raw_data %>%
  distinct(Name, .keep_all = TRUE)

# Now, we check again if there are any duplicated names to see if the remo-
# val was successful
if(n_distinct(raw_data$Name) == nrow(raw_data)) {
  print(paste("No duplicates in patient names"))
} else {
  num_duplicates <- nrow(raw_data)-n_distinct(raw_data$Name)
  print(paste("Number of potential duplicates in patient names:", num_duplicates))
}

## [1] "No duplicates in patient names"

# There are no duplicates in patient names, so the removal was successful

```

## Appendix 2.6: Checking For Negative Values

We can see that all the numeric and integer columns should not have any negative scores, so we need to make sure that is true. The numeric and integer columns are age, billing amount and room number.

```

# The function below selects the chosen numeric and integer values that we
# can observe from the diagnose table above and gives us a summary of how m-
# any negative values are in each of the chosen columns

# Choose the numeric and integer columns
num_cols <- c("Age", "Billing.Amount", "Room.Number")

# Filter to see if there are any negative values
negative_values <- raw_data %>%
  filter(if_any(all_of(num_cols), ~ . < 0))

# Print to see all of them
print(negative_values)

```

##		Name	Age	Gender	Blood.Type	Medical.Condition
## 1		ashley erickson	32	Female	AB-	Cancer
## 2		christopher weiss	49	Female	AB-	Asthma
## 3		ashley warner	60	Male	A+	Hypertension
## 4		jay galloway	74	Female	O+	Asthma
## 5		joshua williamson	72	Female	B-	Diabetes
## 6		scott vazquez	74	Male	B+	Diabetes
## 7		carol anderson	39	Female	B-	Hypertension
## 8	mr.	christopher alvarado	77	Male	AB+	Obesity
## 9		alexandra khan	32	Male	AB+	Arthritis
## 10		joseph cox	23	Male	AB-	Diabetes
## 11		mitchell maldonado	41	Female	A+	Asthma
## 12		julie christian	28	Female	AB+	Diabetes
## 13		alexander gardner	45	Male	AB-	Cancer
## 14		angela allen	83	Female	B-	Diabetes
## 15		alan scott	48	Female	A-	Asthma
## 16		brandon hall	27	Male	B+	Hypertension
## 17		lucas sullivan	31	Female	A-	Obesity
## 18		rebecca cline	39	Female	B+	Asthma
## 19		valerie allen	82	Female	B-	Diabetes
## 20		brian young	85	Male	O-	Cancer
## 21		jennifer cruz	19	Female	A+	Hypertension
## 22		elaine tran	32	Female	O-	Arthritis
## 23		ashley ramsey	53	Male	AB+	Cancer
## 24		john ferrell	58	Female	O-	Hypertension
## 25		christopher hamilton	43	Male	O-	Asthma
## 26		michael castaneda	85	Female	A+	Diabetes
## 27		jason owen	57	Female	B+	Hypertension
## 28		nicole hurst dvm	43	Female	AB+	Obesity
## 29		gabrielle decker	69	Male	AB-	Arthritis
## 30		susan ellis	28	Male	O+	Hypertension
## 31		heather bryant	50	Female	A+	Asthma
## 32		ricardo reynolds	49	Female	O-	Obesity
## 33		whitney cooper	65	Male	AB-	Obesity
## 34		mark stone	33	Male	A+	Diabetes
## 35		frederick moore	83	Male	AB-	Obesity
## 36		stephen chan	71	Male	A+	Cancer
## 37		elizabeth thompson	54	Female	A-	Cancer
## 38		terry wilson	39	Male	O-	Diabetes
## 39		patrick perry	44	Male	B-	Hypertension
## 40		juan osborne	20	Female	AB+	Diabetes
## 41		brenda parrish	59	Male	B+	Diabetes
## 42		melissa diaz	82	Male	A-	Hypertension
## 43	mrs.	michelle clark dvm	31	Female	A-	Asthma
## 44	mrs.	andrea davis phd	25	Female	A+	Hypertension
## 45		daniel drake	67	Female	B+	Hypertension
## 46		alexander martin	46	Male	A+	Hypertension
## 47		walter french	78	Female	A-	Asthma
## 48		craig salazar	19	Female	O+	Diabetes
## 49		molly tapia	78	Female	AB-	Hypertension
## 50		karen williams	31	Female	AB-	Asthma
## 51		calvin campos	26	Male	AB-	Cancer
## 52		donna proctor	32	Female	O+	Obesity
## 53		joseph stevens	75	Female	O+	Obesity

## 54	susan ellison	60	Female	B-	Asthma
## 55	timothy ford	47	Male	AB-	Diabetes
## 56	madeline thomas	32	Female	O-	Arthritis
## 57	bryan rios	85	Female	A-	Cancer
## 58	johnathan smith	18	Female	A-	Arthritis
## 59	veronica kelley	20	Male	O-	Cancer
## 60	john hahn	80	Male	O-	Asthma
## 61	amanda morgan	35	Female	AB+	Hypertension
## 62	dr. michael mckay	67	Male	O+	Cancer
## 63	betty moore md	84	Female	A+	Obesity
## 64	jillian shepherd	84	Female	O-	Diabetes
## 65	susan solis	40	Female	A+	Cancer
## 66	joseph Robbins	71	Male	O-	Hypertension
## 67	erica woods	57	Male	B-	Cancer
## 68	daniel david	32	Male	O+	Obesity
## 69	ryan potter	53	Male	A+	Diabetes
## 70	thomas pratt	57	Female	A+	Obesity
## 71	james byrd	79	Male	O+	Diabetes
## 72	james luna	64	Female	AB-	Cancer
## 73	wanda chase	60	Male	O-	Hypertension
## 74	krista conley	21	Female	O+	Diabetes
## 75	emma savage	57	Male	AB+	Hypertension
## 76	heather rios	81	Female	A-	Asthma
## 77	aaron flowers	42	Male	A+	Asthma
##	Date.of.Admission		Doctor		Hospita
tal					
## 1	2019-11-05		Gerald Hooper		and Johnson Moore, Bra
nch					
## 2	2023-02-16		Kelly Thompson		Hunter-Hug
hes					
## 3	2021-12-21		Andrea Bentley		and Wagner, Lee Kl
ein					
## 4	2021-01-20		Debra Everett		Group Pet
ers					
## 5	2021-03-21		Wendy Ramos		and Huff Reeves, Den
nis					
## 6	2023-04-12		Edward Yates		James
Ltd					
## 7	2020-04-03	Dr. Patrick Hines		Carter Carter, and Patter	
son					
## 8	2022-06-03	Mr. Dean Guzman DDS			Johnson
Inc					
## 9	2022-07-14		Michael Vaughn		Bowen Lopez, and Te
rry					
## 10	2019-10-13		Peter Smith		Inc W
ard					
## 11	2022-06-22		Melissa Woods		Richardson-Benn
ett					
## 12	2019-06-30		Kathryn Ray		Group Mcconn
ell					
## 13	2023-03-23		Tracy Brown		Maldonado Gr
oup					
## 14	2019-08-21		Robert Kennedy		Brown and Randolph, Yo
ung					

## 15 hes	2019-10-23	Jacob Kirby	Lopez-Hug
## 16 ans	2020-01-29	Cindy Smith	Sons and Ev
## 17 ris	2022-04-05	Brenda Burke	Sons and Har
## 18 ton	2020-05-07	Rebecca Miller	Stout Dougherty, and Washingt
## 19 and	2021-02-05	Judy Massey	Hebert, Brown Kim
## 20 LLC	2021-04-13	Dr. Jenna Caldwell	Hampton
## 21 ons	2021-06-09	Samantha Lloyd	Ruiz and S
## 22 son	2022-09-15	Alicia Thompson	Ramirez-Thomp
## 23 ell	2019-05-30	James Smith	Barker-Mitch
## 24 cer	2019-05-20	Randy Calderon	Inc Spen
## 25 oup	2023-01-09	Pamela Obrien	Brock Gr
## 26 and	2023-02-14	Amber Ochoa	Huber, Rodriguez Chapman
## 27 oup	2020-01-27	Matthew Clark	Crawford Gr
## 28 tts	2021-04-19	Leah Russell	Warner-Wa
## 29 hop	2022-11-07	William Krause	Diaz-Bis
## 30 all	2019-12-05	Rebecca Valentine	Anderson-B
## 31 oup	2021-02-26	Morgan Scott	Scott Gr
## 32 les	2019-11-20	Deborah Cox	Wyatt Murphy, and Gonza
## 33 rns	2019-08-20	Amber Walker	and Sons Bu
## 34 rks	2022-10-10	Stephanie Garcia	Cole-Pa
## 35 oup	2019-10-15	Mary Johnson	Miller Gr
## 36 ith	2024-03-05	Allen McGrath	Jones-Sm
## 37 ald	2023-12-11	Jessica Mills DDS Fitzpatrick, Nielsen and Mcdon	
## 38 rt,	2024-02-19	Diana Smith	Medina and Elliott Stewa
## 39 nts	2022-03-26	Miss Sandra Brooks	Foster-Cleme
## 40 son	2021-04-28	Joe Lawson	Garcia-Erick
## 41 ton	2024-04-14	Sherry Brown	and Dennis, McGuire Johns

## 42 nas	2019-06-06	Stephen Levy	Hill-Sali
## 43 oup	2019-08-23	Denise Gilbert	Bruce Gr
## 44 and	2024-03-26	Brian Alvarado	Poole, Poole Mendoza
## 45 Ltd	2020-04-24	Brett Ray	Carr
## 46 now	2021-09-24	Tracy Smith	Carter, and Nguyen S
## 47 sey	2023-08-05	Jeanette Rodriguez	Lyons-Ca
## 48 ler	2022-01-18	Heidi Williams	Group Mil
## 49 son	2024-04-15	Steven Clements	and Rosales, Macdonald Han
## 50 non	2021-10-07	Danny French	Smith-Can
## 51 man	2023-12-15	Eric McCoy	Clements-Bow
## 52 PLC	2023-11-28	Ronald McDonald	Jones
## 53 ons	2023-01-30	John Watson	Schneider and S
## 54 odd	2022-08-28	Vincent Cox	Johnson-T
## 55 ore	2020-08-08	Steven Kirby	Anderson-Mo
## 56 PLC	2022-02-08	Carmen Mann	Moore
## 57 PLC	2022-04-10	Timothy Marshall	Miller
## 58 Ltd	2020-11-24	Sara Watson	Williams
## 59 ing	2021-07-12	Amanda Ramirez	Roberts-K
## 60 son	2024-02-24	Mark Hines	Ltd Wil
## 61 dez	2020-09-27	Gregory Figueroa	Petersen-Hernan
## 62 ark	2019-05-31	Dawn Navarro	Mcconnell and Rios, Cl
## 63 Inc	2022-10-13	Crystal Baker	Aguilar
## 64 oup	2019-12-28	George King	Jones Gr
## 65 ams	2020-04-26	Frances Rosales	Ltd Willi
## 66 PLC	2024-02-10	Stacey Davenport	Vaughn
## 67 ter	2022-10-08	Nathan Gutierrez	LLC Carpen
## 68 oll	2022-05-25	Elizabeth Smith	Estes and Garza, Carr

## 69 and	2022-02-04	Stephen Jensen	Thomas Sons	
## 70 old	2021-03-27	Eduardo Hall	and Wade, Huffman Arn	
## 71 ith	2023-12-25	Adam Pitts	Ayers-Sm	
## 72 ark	2022-09-22	Todd Becker	Juarez-C1	
## 73 on,	2021-03-26	Justin Gibbs	Jimenez and Ross Mas	
## 74 eil	2021-01-03	Elizabeth Anderson	Ltd Mcn	
## 75 LLC	2021-09-05	Maria Johnson	Armstrong	
## 76 and	2021-08-19	Morgan Jackson	Jones, Edwards Jacobs	
## 77 ard	2020-06-03	Ronald Patton	Brown, Santos and How	
## Insurance.Provider Billing.Amount Room.Number Admission.Type Discharge.Date				
## 1 9-11-23	Aetna	-502.50781	376	Urgent
## 2 3-03-09	Aetna	-1018.24537	204	Elective
## 3 2-01-11	Aetna	-306.36493	426	Elective
## 4 1-02-09	Blue Cross	-109.09712	381	Emergency
## 5 1-04-17	Blue Cross	-576.72791	369	Urgent
## 6 3-05-03	Medicare	-135.98600	445	Elective
## 7 0-04-10	Blue Cross	-370.98367	203	Elective
## 8 2-06-13	Blue Cross	-1310.27289	257	Elective
## 9 2-07-19	Aetna	-692.40882	372	Elective
## 10 9-10-25	Blue Cross	-353.86519	271	Elective
## 11 2-06-30	Medicare	-378.96078	414	Urgent
## 12 9-07-09	Cigna	-367.20396	387	Urgent
## 13 3-04-09	Cigna	-198.28381	329	Elective
## 14 9-08-24	Cigna	-43.09852	189	Urgent
## 15 9-11-18	Medicare	-857.12593	309	Urgent
## 16 0-02-01	Aetna	-155.08202	368	Urgent
## 17 2-04-11	UnitedHealthcare	-211.72385	496	Emergency

## 18 0-05-17	Cigna	-656.15307	195	Emergency	202
## 19 1-02-25	Medicare	-227.99506	118	Urgent	202
## 20 1-04-20	Cigna	-147.07220	133	Emergency	202
## 21 1-06-28	Blue Cross	-124.75643	207	Urgent	202
## 22 2-10-12	UnitedHealthcare	-75.81945	482	Elective	202
## 23 9-06-09	Blue Cross	-135.71907	235	Elective	201
## 24 9-05-27	Medicare	-308.58427	394	Emergency	201
## 25 3-01-25	Medicare	-214.60542	398	Elective	202
## 26 3-03-09	Medicare	-224.63242	179	Emergency	202
## 27 0-02-16	Blue Cross	-100.60361	354	Elective	202
## 28 1-04-23	UnitedHealthcare	-211.28374	213	Urgent	202
## 29 2-11-30	Cigna	-676.85250	354	Urgent	202
## 30 0-01-02	UnitedHealthcare	-130.86753	451	Elective	202
## 31 1-03-07	UnitedHealthcare	-233.93073	289	Emergency	202
## 32 9-12-20	Aetna	-124.64904	222	Emergency	201
## 33 9-09-02	UnitedHealthcare	-577.72911	357	Emergency	201
## 34 2-10-23	Cigna	-53.83209	457	Urgent	202
## 35 9-11-12	Cigna	-532.76112	132	Urgent	201
## 36 4-03-12	Cigna	-599.26536	119	Urgent	202
## 37 3-12-13	Aetna	-887.02422	402	Elective	202
## 38 4-03-20	Aetna	-1316.61858	491	Emergency	202
## 39 2-03-29	UnitedHealthcare	-75.62019	340	Urgent	202
## 40 1-05-20	Blue Cross	-317.63292	324	Emergency	202
## 41 4-04-22	UnitedHealthcare	-90.07890	221	Emergency	202
## 42 9-07-06	Aetna	-23.86673	453	Emergency	201
## 43 9-09-14	Cigna	-952.83119	136	Urgent	201
## 44 4-04-22	Aetna	-492.17839	372	Emergency	202

## 45 0-04-26	Aetna	-591.91742	426	Elective	202
## 46 1-09-26	UnitedHealthcare	-38.96615	401	Emergency	202
## 47 3-08-30	UnitedHealthcare	-136.16469	347	Emergency	202
## 48 2-01-26	UnitedHealthcare	-37.26775	409	Elective	202
## 49 4-05-05	UnitedHealthcare	-230.57928	211	Elective	202
## 50 1-10-15	Cigna	-407.80082	494	Emergency	202
## 51 4-01-13	Medicare	-1277.64534	339	Elective	202
## 52 3-12-28	Aetna	-786.62472	435	Elective	202
## 53 3-02-13	Cigna	-967.59471	476	Elective	202
## 54 2-09-08	Medicare	-416.91482	365	Urgent	202
## 55 0-08-23	UnitedHealthcare	-97.85312	169	Elective	202
## 56 2-02-25	Aetna	-964.79862	122	Urgent	202
## 57 2-04-11	Cigna	-311.75563	471	Urgent	202
## 58 0-12-03	Cigna	-85.47571	349	Urgent	202
## 59 1-07-28	Medicare	-808.46506	221	Emergency	202
## 60 4-03-14	Cigna	-1520.42055	403	Elective	202
## 61 0-10-20	Medicare	-652.18137	144	Elective	202
## 62 9-06-12	UnitedHealthcare	-199.66379	122	Urgent	201
## 63 2-10-26	Blue Cross	-529.96301	409	Urgent	202
## 64 0-01-07	UnitedHealthcare	-279.84241	142	Emergency	202
## 65 0-04-29	Cigna	-226.38112	132	Emergency	202
## 66 4-02-29	UnitedHealthcare	-1428.84394	205	Urgent	202
## 67 2-10-14	Cigna	-1049.01234	146	Urgent	202
## 68 2-06-22	Aetna	-36.21727	167	Emergency	202
## 69 2-03-01	Medicare	-614.94559	223	Urgent	202
## 70 1-04-22	Medicare	-228.54685	496	Emergency	202
## 71 3-12-29	Cigna	-483.70889	446	Urgent	202

## 72	Aetna	-2008.49214	162	Urgent	202
2-10-20					
## 73	Medicare	-820.16335	357	Emergency	202
1-04-12					
## 74	Aetna	-289.80162	340	Urgent	202
1-01-14					
## 75	Blue Cross	-1660.00937	193	Elective	202
1-10-02					
## 76	Blue Cross	-378.74681	249	Urgent	202
1-08-23					
## 77	Medicare	-378.69641	115	Elective	202
0-06-17					
## Medication Test.Results					
## 1	Penicillin	Normal			
## 2	Penicillin	Inconclusive			
## 3	Ibuprofen	Normal			
## 4	Ibuprofen	Abnormal			
## 5	Aspirin	Abnormal			
## 6	Ibuprofen	Abnormal			
## 7	Ibuprofen	Abnormal			
## 8	Paracetamol	Inconclusive			
## 9	Lipitor	Abnormal			
## 10	Lipitor	Inconclusive			
## 11	Lipitor	Inconclusive			
## 12	Lipitor	Normal			
## 13	Aspirin	Abnormal			
## 14	Penicillin	Abnormal			
## 15	Paracetamol	Abnormal			
## 16	Aspirin	Abnormal			
## 17	Aspirin	Abnormal			
## 18	Ibuprofen	Abnormal			
## 19	Aspirin	Inconclusive			
## 20	Aspirin	Abnormal			
## 21	Aspirin	Normal			
## 22	Lipitor	Abnormal			
## 23	Penicillin	Inconclusive			
## 24	Paracetamol	Inconclusive			
## 25	Ibuprofen	Inconclusive			
## 26	Penicillin	Inconclusive			
## 27	Paracetamol	Inconclusive			
## 28	Paracetamol	Inconclusive			
## 29	Penicillin	Normal			
## 30	Lipitor	Inconclusive			
## 31	Penicillin	Inconclusive			
## 32	Lipitor	Normal			
## 33	Aspirin	Normal			
## 34	Aspirin	Inconclusive			
## 35	Aspirin	Normal			
## 36	Penicillin	Abnormal			
## 37	Ibuprofen	Normal			
## 38	Paracetamol	Inconclusive			
## 39	Aspirin	Inconclusive			
## 40	Ibuprofen	Inconclusive			
## 41	Penicillin	Abnormal			

```

## 42 Ibuprofen      Abnormal
## 43 Penicillin     Normal
## 44 Aspirin        Inconclusive
## 45 Lipitor         Abnormal
## 46 Lipitor         Inconclusive
## 47 Aspirin         Abnormal
## 48 Aspirin         Abnormal
## 49 Penicillin     Inconclusive
## 50 Aspirin         Normal
## 51 Ibuprofen      Normal
## 52 Paracetamol    Abnormal
## 53 Ibuprofen      Normal
## 54 Aspirin         Abnormal
## 55 Paracetamol    Inconclusive
## 56 Paracetamol    Normal
## 57 Paracetamol    Abnormal
## 58 Paracetamol    Normal
## 59 Paracetamol    Inconclusive
## 60 Lipitor          Abnormal
## 61 Penicillin      Abnormal
## 62 Ibuprofen      Abnormal
## 63 Aspirin         Inconclusive
## 64 Penicillin      Inconclusive
## 65 Lipitor          Inconclusive
## 66 Ibuprofen      Abnormal
## 67 Lipitor          Normal
## 68 Penicillin      Inconclusive
## 69 Penicillin      Abnormal
## 70 Paracetamol    Inconclusive
## 71 Penicillin      Abnormal
## 72 Ibuprofen      Abnormal
## 73 Lipitor          Inconclusive
## 74 Lipitor          Inconclusive
## 75 Penicillin      Normal
## 76 Lipitor          Abnormal
## 77 Paracetamol    Inconclusive

# The output shows that there are negative values, so there we will remove them

raw_data <- raw_data %>%
  filter(if_all(all_of(num_cols), ~ . >= 0))

# Run negative values again to check if removal was successful
negative_values <- raw_data %>%
  filter(if_any(all_of(num_cols), ~ . < 0))

print(negative_values)

## [1] Name           Age            Gender          Blood.Type
## [5] Medical.Condition Date.of.Admission Doctor          Hospital
## [9] Insurance.Provider Billing.Amount   Room.Number    Admission

```

```

.Type
## [13] Discharge.Date      Medication          Test.Results
## <0 rows> (or 0-length row.names)

# Since there are no rows printed, the removal was successful and there are no more negative values where there should not be.

```

## Appendix 2.7: Check The Type Of The Data

```

diagnosis <- diagnose(raw_data)
diagnosis

## # A tibble: 15 × 6
##   variables     types missing_count missing_percent unique_count un
##   <chr>        <chr>    <int>           <dbl>       <int>
## 1 Name         char...     0             0            40158
## 2 Age          inte...     0             0            68
## 3 Gender       char...     0             0            2
## 4 Blood.Type   char...     0             0            8
## 5 Medical.Condition char...     0             0            6
## 6 Date.of.Admission char...     0             0            1827
## 7 Doctor       char...     0             0            33502
## 8 Hospital     char...     0             0            32734
## 9 Insurance.Provider char...     0             0            5
## 10 Billing.Amount nume...     0             0            40158
## 11 Room.Number  inte...     0             0            400
## 12 Admission.Type char...     0             0            3
## 13 Discharge.Date  char...     0             0            1856
## 14 Medication    char...     0             0            5
## 15 Test.Results  char...     0             0            3
## 0.000125
## 0.000747

# From here, we can tell that there are types of data that need to be changed, so we assign the correct types to each variables

raw_data$Gender <- as.factor(raw_data$Gender)
raw_data$Blood.Type <- as.factor(raw_data$Blood.Type)
raw_data$Medical.Condition <- as.factor(raw_data$Medical.Condition)
raw_data>Date.of.Admission <- as.Date(raw_data>Date.of.Admission)

```

```

raw_data$Doctor <- as.factor(raw_data$Doctor)
raw_data$Hospital <- as.factor(raw_data$Hospital)
raw_data$Insurance.Provider <- as.factor(raw_data$Insurance.Provider)
raw_data$Admission.Type <- as.factor(raw_data$Admission.Type)
raw_data$Discharge.Date <- as.Date(raw_data$Discharge.Date)
raw_data$Medication <- as.factor(raw_data$Medication)
raw_data$Test.Results <- as.factor(raw_data$Test.Results)

diagnosis <- diagnose(raw_data)
diagnosis

## # A tibble: 15 × 6
##   variables      types missing_count missing_percent unique_count un
##   <chr>        <chr>          <int>             <dbl>          <int>
## 1 Name          char...          0                 0       40158
## 2 Age           inte...          0                 0       68
## 3 Gender         fact...          0                 0       2
## 4 Blood.Type    fact...          0                 0       8
## 5 Medical.Conditi... fact...          0                 0       6
## 6 Date.of.Admissi... Date            0                 0      1827
## 7 Doctor         fact...          0                 0      33502
## 8 Hospital        fact...          0                 0      32734
## 9 Insurance.Provi... fact...          0                 0       5
## 10 Billing.Amount nume...          0                 0       40158
## 11 Room.Number   inte...          0                 0       400
## 12 Admission.Type fact...          0                 0       3
## 13 Discharge.Date Date            0                 0      1856
## 14 Medication     fact...          0                 0       5
## 15 Test.Results   fact...          0                 0       3
## 0.0000747

```

## Appendix 2.8: Removing Unnecessary Columns

There are some columns that were only used to clean the data, but are not relevant to the analysis, so we will remove them.

Here, we will remove the “Name”, “Doctor”, “Date of Admission”, “Hospital” and “Discharge Date” columns as they were only used to identify duplicates, but will not affect analysis at all.

```
# Removing "Name", "Doctor", "Date of Admission", "Hospital" and "Discharge Date" columns
raw_data <- raw_data[, -c(1, 6:8, 13)]

# Run diagnosis to see if "Name" column has been removed
diagnosis <- diagnose(raw_data)
diagnosis

## # A tibble: 10 × 6
##   variables     types missing_count missing_percent unique_count un
##   <chr>        <chr>      <int>             <dbl>       <int>
##   <dbl>
## 1 Age          integer     0                 0           68
## 2 Gender        factor      0                 0           2
## 3 Blood.Type    factor      0                 0           8
## 4 Medical.Conditi... factor      0                 0           6
## 5 Insurance.Provi... factor      0                 0           5
## 6 Billing.Amount numeric     0                 0         40158
## 7 Room.Number   integer     0                 0           400
## 8 Admission.Type factor      0                 0           3
## 9 Medication    factor      0                 0           5
## 10 Test.Results  factor      0                 0           3
## 11
```

## Appendix 2.9: Renaming Checked And Cleaned Data Set

Hence, we can rename “raw\_data” as “health”

```
health <- raw_data # Rename the "raw_data" to "health"

# Use diagnosis() function to check whether renaming is successful. If renaming is successful, it will show the same diagnosis table from above
diagnosis <- diagnose(health)
diagnosis

## # A tibble: 10 × 6
##   variables     types missing_count missing_percent unique_count un
##   <chr>        <chr>      <int>             <dbl>       <int>
##   <dbl>
## 1 Age          integer     0                 0           68
```

```

0.00169
## 2 Gender      fact...      0      0      2
0.0000498
## 3 Blood.Type   fact...      0      0      8
0.000199
## 4 Medical.Conditi... fact...      0      0      6
0.000149
## 5 Insurance.Provi... fact...      0      0      5
0.000125
## 6 Billing.Amount  nume...      0      0  40158
1
## 7 Room.Number    inte...      0      0      400
0.00996
## 8 Admission.Type  fact...      0      0      3
0.0000747
## 9 Medication     fact...      0      0      5
0.000125
## 10 Test.Results   fact...      0      0      3
0.0000747

```

## Appendix 3: Exploratory Data Analysis (EDA)

### Appendix 3.1: Loading Necessary Packages

```

library(dplyr) # For data manipulation
library(tidyr) # To tidy data
library(ggplot2) # Load package ggplot2 for function ggPlot to make plots
library(scales) # This package allows us to transform any scientific notation on axes to be displayed as normal numbers

##
## Attaching package: 'scales'

## The following object is masked from 'package:purrr':
##
##     discard

## The following object is masked from 'package:readr':
##
##     col_factor

library(janitor) # To get the proportion of each discrete variable

##
## Attaching package: 'janitor'

## The following objects are masked from 'package:stats':
##
##     chisq.test, fisher.test

```

### Appendix 3.2: Basic Statistics and Key Points Of Data

Using the summary function, we can better understand the key points and basic statistics in each variable before going into further analysis.

### Appendix 3.2.1: Basic Summary

```
summary(health) # Shows the type and general key characteristics of each variable, though it is only more useful for the continuous distributions as it can show the key points in the continuous data

##      Age          Gender        Blood.Type       Medical.Condition
## Min.   :18.0    Female:20084    AB+: 5093    Arthritis   :6816
## 1st Qu.:35.0   Male  :20074     B+ : 5059    Asthma     :6603
## Median  :52.0                    B- : 5020    Cancer     :6675
## Mean    :51.7                    A- : 5014    Diabetes   :6779
## 3rd Qu.:69.0                    A+ : 5013    Hypertension:6661
## Max.    :85.0                    O- : 5006    Obesity    :6624
##                               (Other):9953
##      Insurance.Provider Billing.Amount      Room.Number
## Aetna           :7929      Min.   : 9.239   Min.   :101.0
## Blue Cross     :8066      1st Qu.:13292.171 1st Qu.:202.0
## Cigna          :8104      Median  :25595.331 Median  :302.0
## Medicare       :8103      Mean    :25609.534 Mean   :300.9
## UnitedHealthcare:7956    3rd Qu.:37862.835 3rd Qu.:401.0
##                               Max.   :52764.277 Max.   :500.0
##
##      Admission.Type      Medication        Test.Results
## Elective  :13538    Aspirin   :8038    Abnormal   :13457
## Emergency:13099  Ibuprofen:8016    Inconclusive:13278
## Urgent    :13521    Lipitor   :8141    Normal    :13423
##                               Paracetamol:7963
##                               Penicillin:8000
##
##
```

Analysis: - Age: Has a range from 18 to 85 with an average of 51.7 years old - Billing amount: Has a range of 9.24 to 52764.28 with similar mean and median - Room number: Has a range from 101 to 500

### Appendix 3.2.2: Discrete Data

We use the tabyl() function to create frequency tables (counts and proportions) of each variable in a data frame. This is for use to get an idea of how the discrete data is structured.

```
# First, we look for which column is a factor column. This will tell us which variables are discrete
sapply(health, is.factor)

##      Age          Gender        Blood.Type       Medical.Condition
## FALSE         TRUE          TRUE          TRUE
## Insurance.Provider Billing.Amount      Room.Number      Admission.Type
## TRUE          FALSE         FALSE         FALSE
## Medication      Test.Results
## TRUE          TRUE
```

```

# Then, we can use the tabyl() function to create the frequency tables
tabyl(health$Gender)

##  health$Gender      n    percent
##            Female 20084 0.5001245
##            Male  20074 0.4998755

tabyl(health$Blood.Type)

##  health$Blood.Type      n    percent
##            A- 5014 0.1248568
##            A+ 5013 0.1248319
##            AB- 4989 0.1242343
##            AB+ 5093 0.1268240
##            B- 5020 0.1250062
##            B+ 5059 0.1259774
##            O- 5006 0.1246576
##            O+ 4964 0.1236117

tabyl(health$Medical.Condition)

##  health$Medical.Condition      n    percent
##            Arthritis 6816 0.1697296
##            Asthma   6603 0.1644255
##            Cancer   6675 0.1662184
##            Diabetes 6779 0.1688082
##            Hypertension 6661 0.1658698
##            Obesity  6624 0.1649485

tabyl(health$Insurance.Provider)

##  health$Insurance.Provider      n    percent
##            Aetna 7929 0.1974451
##            Blue Cross 8066 0.2008566
##            Cigna 8104 0.2018029
##            Medicare 8103 0.2017780
##            UnitedHealthcare 7956 0.1981174

tabyl(health$Admission.Type)

##  health$Admission.Type      n    percent
##            Elective 13538 0.3371184
##            Emergency 13099 0.3261866
##            Urgent   13521 0.3366951

tabyl(health$Medication)

##  health$Medication      n    percent
##            Aspirin 8038 0.2001594
##            Ibuprofen 8016 0.1996115
##            Lipitor  8141 0.2027242
##            Paracetamol 7963 0.1982917
##            Penicillin 8000 0.1992131

tabyl(health$Test.Results)

```

```

##  health$Test.Results      n   percent
##                Abnormal 13457 0.3351013
##                Inconclusive 13278 0.3306440
##                Normal 13423 0.3342547

```

Analysis: Overall, looks to be quite evenly spread out for all variables. This can be confirmed later on with pie charts and box plots

### Appendix 3.2.3: Correlation

Further, we check for correlation among numeric predictors.

```

cor(health[, sapply(health, is.numeric)], use = "complete.obs")

##                               Age Billing.Amount Room.Number
## Age             1.0000000000  0.0002776157  0.002706352
## Billing.Amount 0.0002776157  1.0000000000 -0.002252990
## Room.Number    0.0027063515 -0.0022529897  1.000000000

```

From what we can see there isn't any significant correlation between any of the numeric variable as none of the correlation coefficients other than the diagonal (which is the correlation the variable itself) is more than 0.5.

Hence, when doing the variable selection in later sections, we will not use any interactions between the variables.

## Appendix 3.3: EDA - Boxplots For Outliers

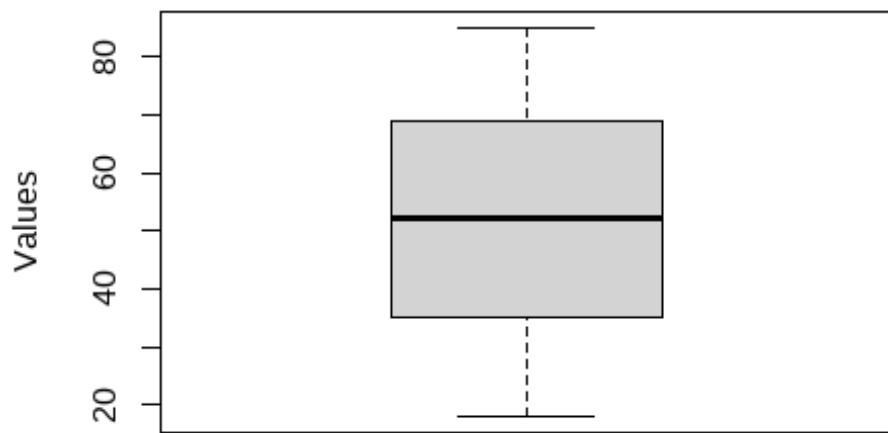
By plotting boxplots for the continuously distributed variables, we can see if there are any outliers present in the variable. Then, we will carry out further analysis to determine whether the outliers need to be removed from the variable to prevent it from affecting the further analysis and predictions.

```

# Loop to make boxplot for each variable
for(var in colnames(health)) {
  # Choose only numeric variables to Look at the outliers as an overview
  if(is.numeric(health[[var]])) {
    # If column is numeric, then create boxplot
    var_box <- boxplot(health[[var]], main=paste("Boxplot of" , var), ylab
    ="Values") + theme_minimal()
    var_box # Show the boxplot created
  } else {
    # If column is not numeric, then says its not applicable
    print(paste("Boxplot is not applicable for", var))
  }
}

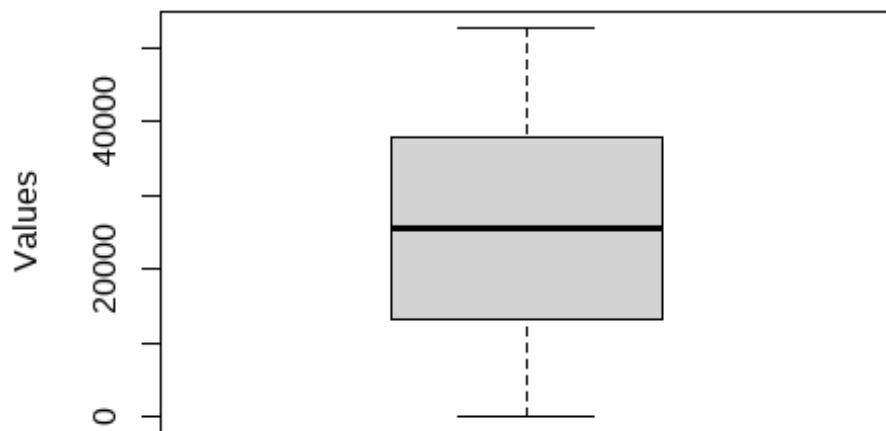
```

### Boxplot of Age

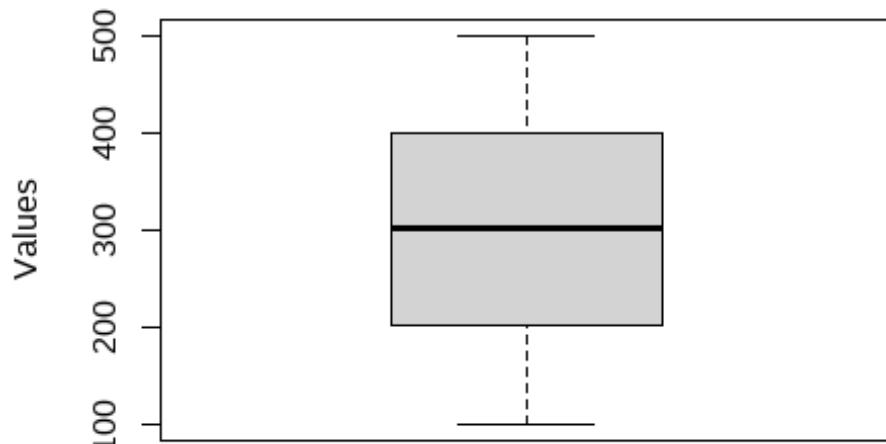


```
## [1] "Boxplot is not applicable for Gender"  
## [1] "Boxplot is not applicable for Blood.Type"  
## [1] "Boxplot is not applicable for Medical.Condition"  
## [1] "Boxplot is not applicable for Insurance.Provider"
```

### Boxplot of Billing.Amount



### Boxplot of Room.Number



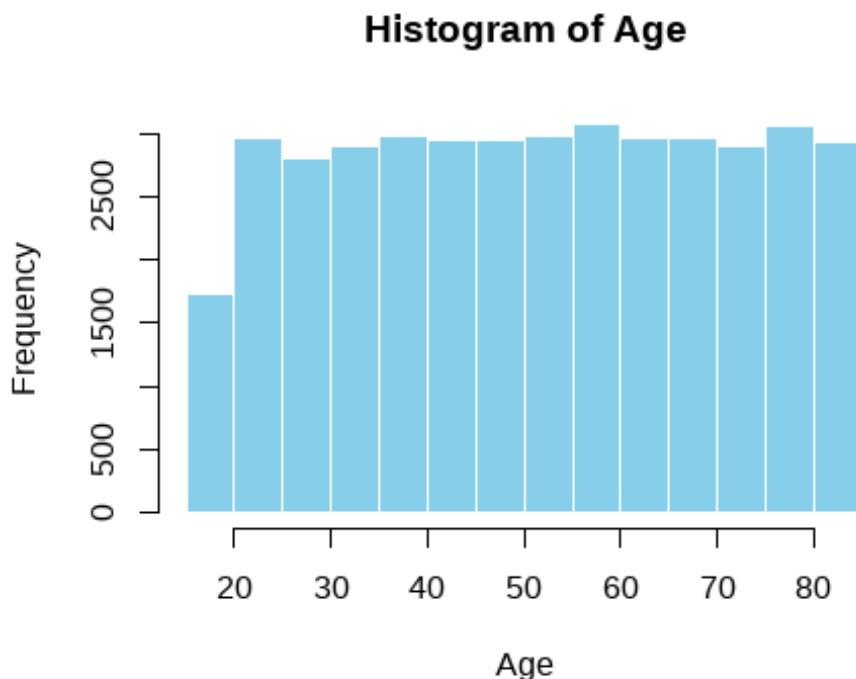
```
## [1] "Boxplot is not applicable for Admission.Type"  
## [1] "Boxplot is not applicable for Medication"  
## [1] "Boxplot is not applicable for Test.Results"
```

Analysis: Results show that there are no outliers. Hence, there is no need to check whether any outliers need to be removed.

## Appendix 3.4: EDA - Histogram

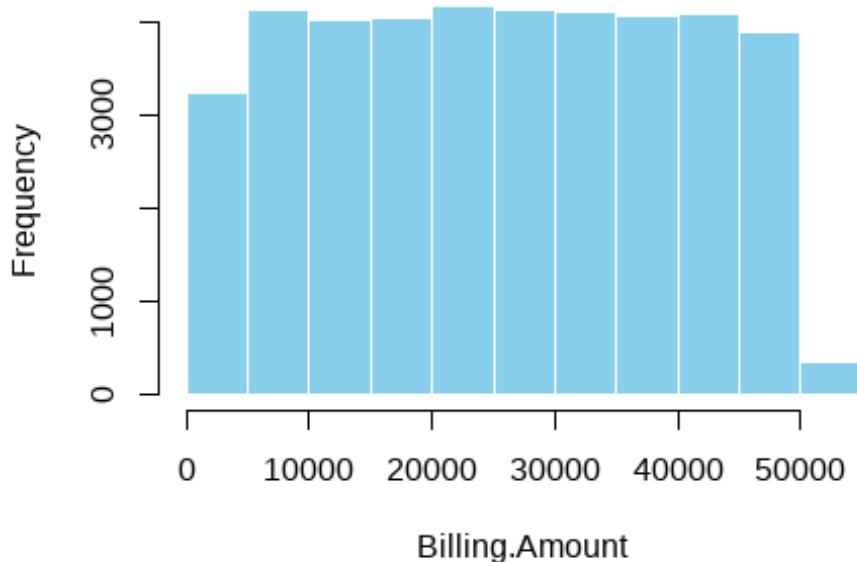
We plot histograms for continuous numeric variables to get an idea of how the variable is distributed.

```
for (var in colnames(health)){
  # Check if the variable is numeric
  if (is.numeric(health[[var]])){
    # Create histogram for each numeric variable
    hist(health[[var]], main = paste("Histogram of", var), xlab = var, col =
    "skyblue", border = "white")
  } else {
    # If column is not numeric, then says its not applicable
    print(paste("Histogram is not applicable for", var))
  }
}
```

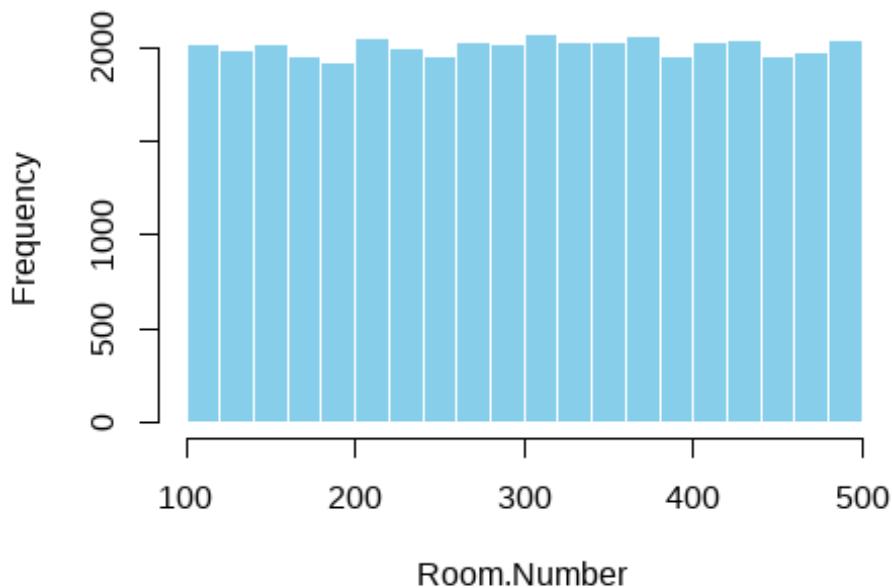


```
## [1] "Histogram is not applicable for Gender"
## [1] "Histogram is not applicable for Blood.Type"
## [1] "Histogram is not applicable for Medical.Condition"
## [1] "Histogram is not applicable for Insurance.Provider"
```

### Histogram of Billing.Amount



### Histogram of Room.Number



```
## [1] "Histogram is not applicable for Admission.Type"  
## [1] "Histogram is not applicable for Medication"  
## [1] "Histogram is not applicable for Test.Results"
```

Analysis: - Age: Looks to be negatively skewed - Billing amount: Looks to be positively skewed - Room number: Has the form of a uniform distribution

## Appendix 3.5: EDA - Pie Charts For All Variables

We use pie charts to see the percentage of each category in the variable to better understand how the data of each variable is proportioned.

### Appendix 3.5.1: Creating Function For Making Pie Charts

Here, we are going to create a unique function to quickly make pie charts. Then, we will loop our unique functions in Section 3.5.2 to each of the variables so that it will create the pie charts easily.

```
pie_chart <- function(data, var) {  
  # Select the variable from the dataset  
  selected_var <- data[[var]]  
  
  # If NA exists, convert it to be an "Unknown" so that it can be classed  
  # as a category as well  
  selected_var[is.na(selected_var)] <- "Unknown"  
  
  # Count the number/frequency  
  count <- table(selected_var)  
  
  # Calculate percentages to 1 decimal place  
  percentages <- round(100*count/sum(count), 1)  
  
  # Adjust the plot margins for the pie chart  
  par(mar=c(1,3,3,1))  
  
  # Create the pie chart  
  draw_chart <- pie(count,  
                    col=rainbow(length(count)), # Auto-assign colours to each of the proportions  
                    main=paste("Pie Chart for", var), # Name of pie chart  
                    labels=paste(percentages, "%"), # Show it as percentage  
                    radius=0.85,  
                    cex=1.1, # Adjust the size of the chart  
                    xpd=TRUE) # Allows legend to be drawn outside the plot region  
  
  # Add legend to see what colour means  
  legend("topleft",  
        legend=names(count),  
        fill=rainbow(length(count)),  
        title="Categories",  
        cex=0.9, # Adjust the size of the text  
        xpd=TRUE)  
}
```

### Appendix 3.5.2: Making Pie Charts

```
# Here, we ignore "Doctor" and "Hospital" column as they have too many factor levels  
for (var in colnames(health)) {  
  if (is.factor(health[[var]])) {  
    pie_chart(health, var)
```

```

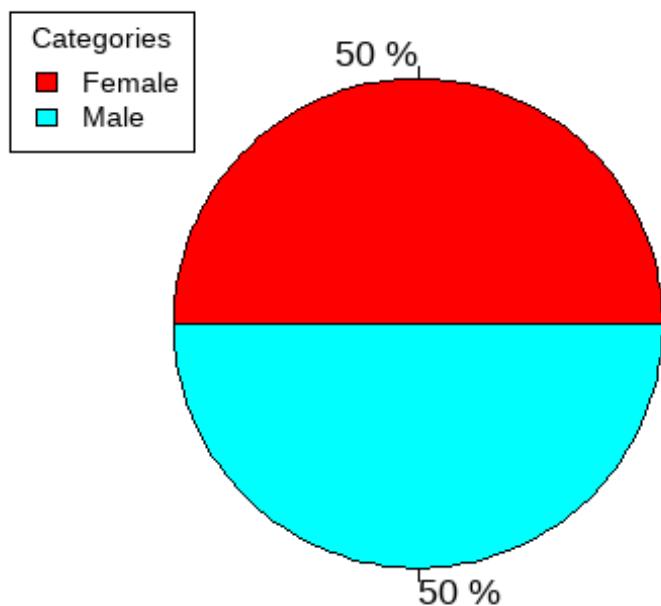
} else {
# If column is not a factor, then says its not applicable
print(paste("Pie chart is not applicable for", var))
}
}

## [1] "Pie chart is not applicable for Age"

## Warning in `[<-.factor`(`*tmp*`, is.na(selected_var), value = "Unknown"
):
## invalid factor level, NA generated

```

**Pie Chart for Gender**

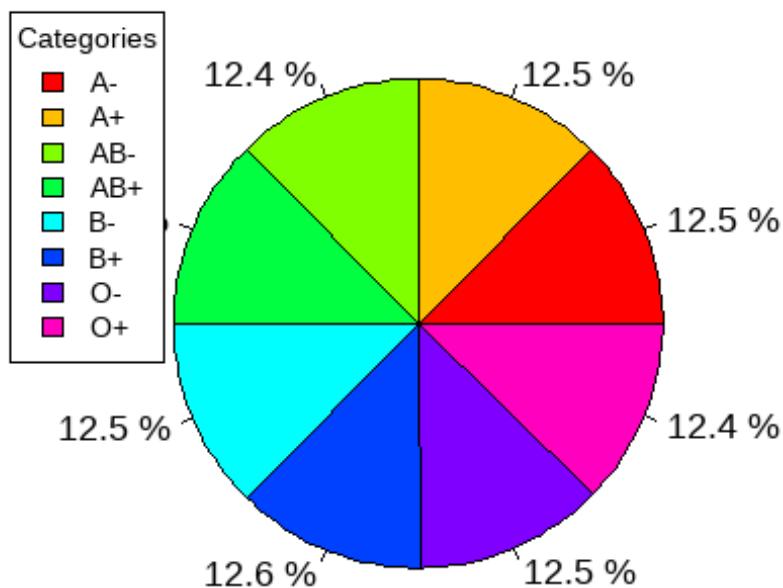


```

## Warning in `[<-.factor`(`*tmp*`, is.na(selected_var), value = "Unknown"
):
## invalid factor level, NA generated

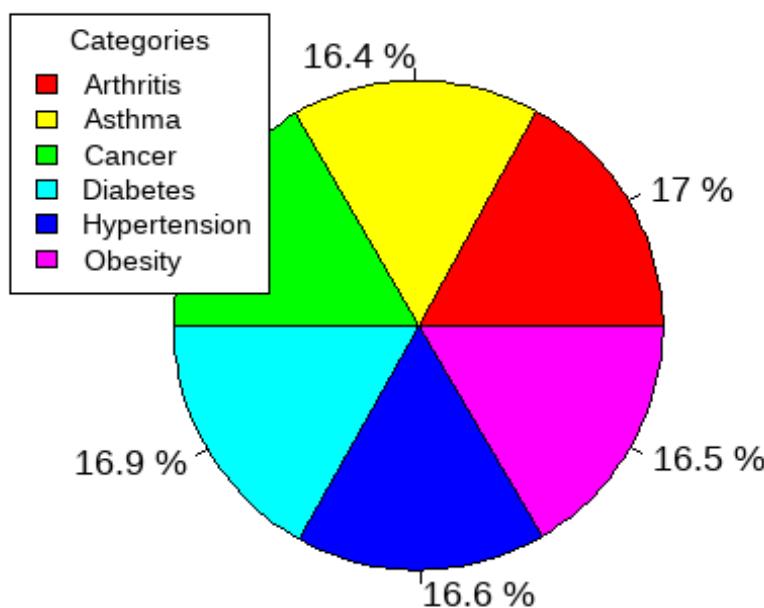
```

### Pie Chart for Blood.Type



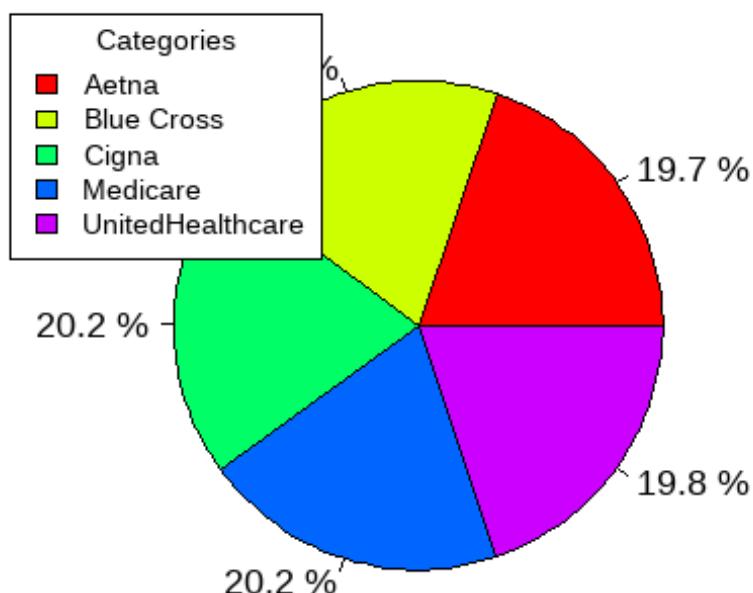
```
## Warning in `[<-.factor`(`*tmp*`, is.na(selected_var), value = "Unknown"):  
## invalid factor level, NA generated
```

### Pie Chart for Medical.Condition



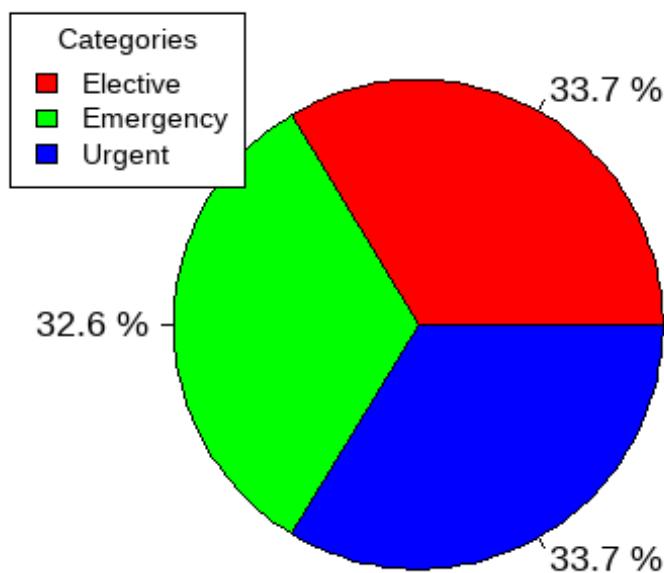
```
## Warning in `[<-.factor`(`*tmp*`, is.na(selected_var), value = "Unknown"):  
## invalid factor level, NA generated
```

### Pie Chart for Insurance.Provider



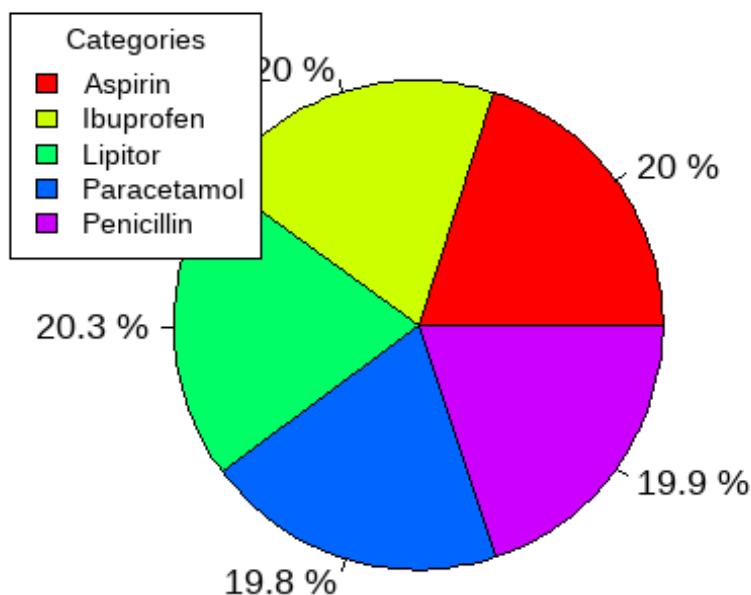
```
## [1] "Pie chart is not applicable for Billing.Amount"  
## [1] "Pie chart is not applicable for Room.Number"  
  
## Warning in `<-.factor`(`*tmp*`, is.na(selected_var), value = "Unknown"  
## :  
## invalid factor level, NA generated
```

### Pie Chart for Admission.Type



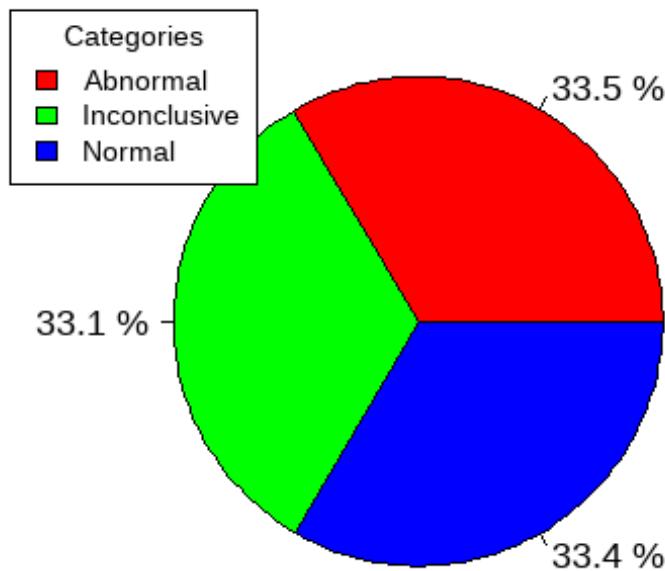
```
## Warning in `[<-.factor`(`*tmp*`, is.na(selected_var), value = "Unknown"):  
## invalid factor level, NA generated
```

**Pie Chart for Medication**



```
## Warning in `[<-.factor`(`*tmp*`, is.na(selected_var), value = "Unknown"):  
## invalid factor level, NA generated
```

## Pie Chart for Test.Results



Analysis: -

Gender: Split evenly between male and female - Blood type, medical condition, insurance provider, admission type, medication, test results: Almost split evenly in its own variables

### Appendix 3.6: EDA - Bar Charts For All Variables

We use bar charts to see the count of each category in the variable to better understand how the data of each variable is spread.

#### Appendix 3.6.1: Creating Function For Making Bar Chart

Here, we are going to create a unique function to quickly make the bar charts. Then, we will loop our unique functions in Section 3.6.2 to each of the variables so that it will create the bar charts easily.

```
bar_chart <- function(data, var) {  
  # Select the variable from the dataset  
  selected_var <- data[[var]]  
  
  # If NA exists, convert it to be an "Unknown" so that it can be classed  
  # as a category as well  
  selected_var[is.na(selected_var)] <- "Unknown"  
  
  # Count the number/frequency  
  count <- as.data.frame(table(selected_var))  
  
  # Rename the columns for ggplot2 compatibility  
  colnames(count) <- c("category", "count")  
  
  # Create bar chart  
  bar <- ggplot(count, aes(x = category, y = count, fill = category)) +
```

```

    geom_bar(stat = "identity") + # Use actual counts
    geom_text(aes(label = count), vjust = -0.5, size = 4) + # Add count labels above bars
    labs(title = paste("Bar Chart - ", var), x = "Categories", y = "Frequency") + # Labels
    scale_fill_manual(values = rainbow(length(count$category))) + # Manual color palette
    theme_minimal() + # Clean theme
    theme(axis.text.x=element_text(angle = 45, hjust = 1)) + # Rotate x-axis labels by 45 degrees
    theme(plot.title=element_text(size=14, face="bold"), # Adjust plot title size,
          axis.text.x = element_text(size = 11),      # Adjust x-axis label size,
          legend.position="none")

# Print the bar chart
print(bar)
}

```

### Appendix 3.6.2: Making Bar Charts

```

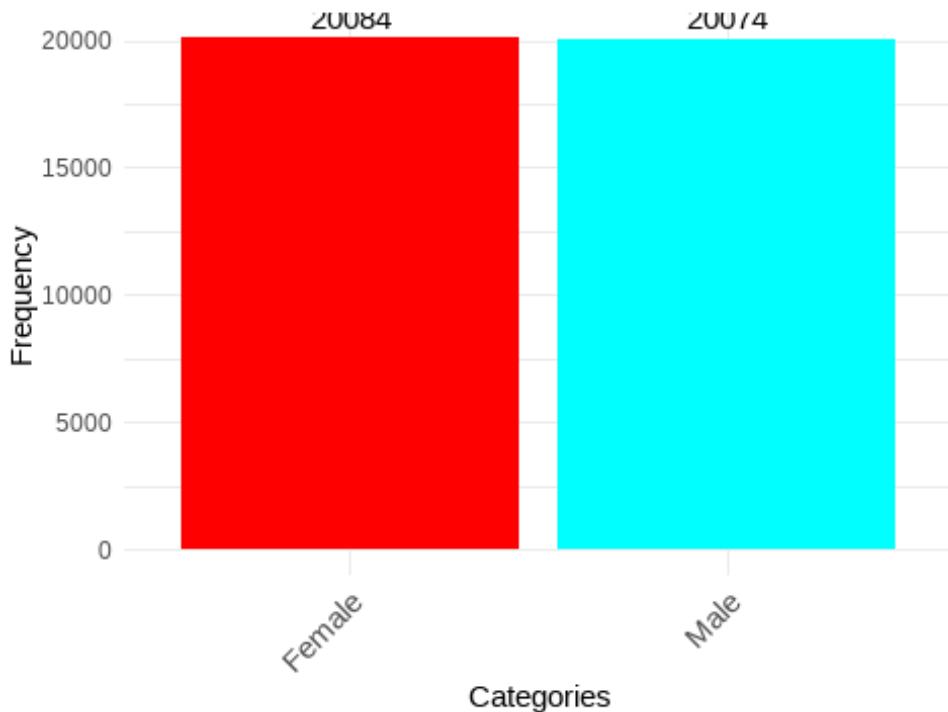
# Here, we there have too many factor Levels
for (var in colnames(health)) {
  if (is.factor(health[[var]])) {
    bar_chart(health, var)
  } else {
    # If column is not a factor, then says its not applicable
    print(paste("Bar chart is not applicable for", var))
  }
}

## [1] "Bar chart is not applicable for Age"

## Warning in `[<-.factor`(`*tmp*`, is.na(selected_var), value = "Unknown"):
## invalid factor level, NA generated

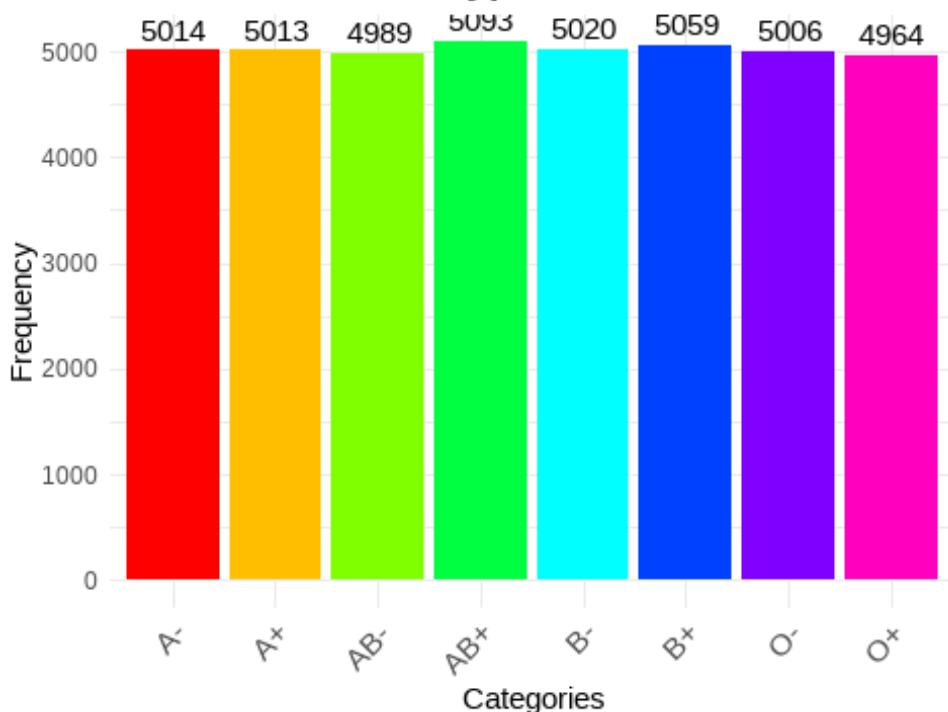
```

### Bar Chart - Gender



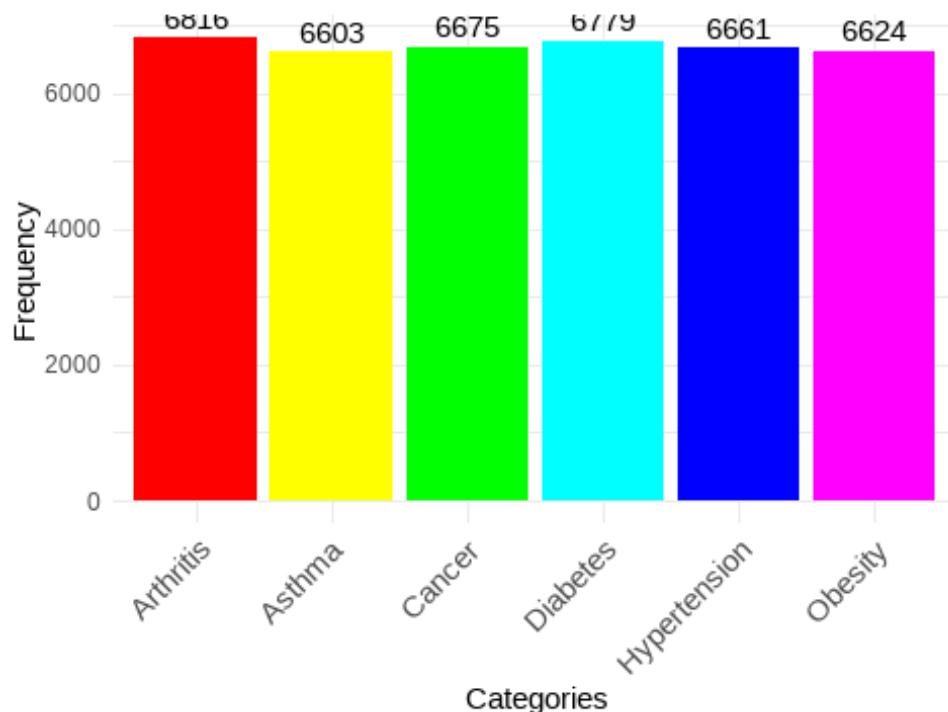
```
## Warning in `[<-.factor`(`*tmp*`, is.na(selected_var), value = "Unknown"):  
## invalid factor level, NA generated
```

### Bar Chart - Blood.Type



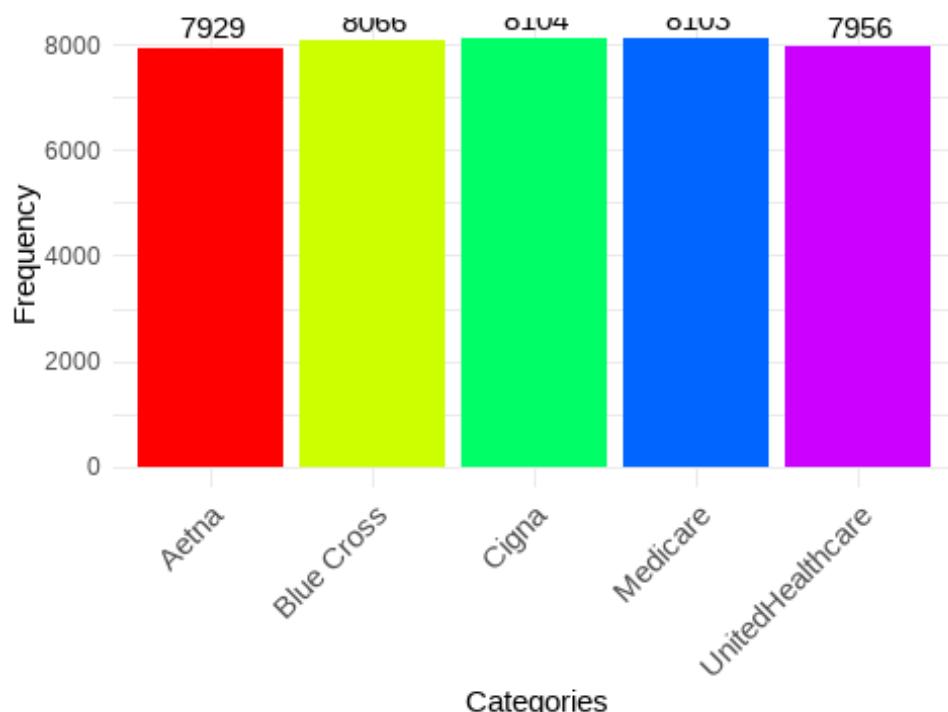
```
## Warning in `[<-.factor`(`*tmp*`, is.na(selected_var), value = "Unknown"):  
## invalid factor level, NA generated
```

### Bar Chart - Medical.Condition



```
## Warning in `[<-.factor`(`*tmp*`, is.na(selected_var), value = "Unknown"):  
## invalid factor level, NA generated
```

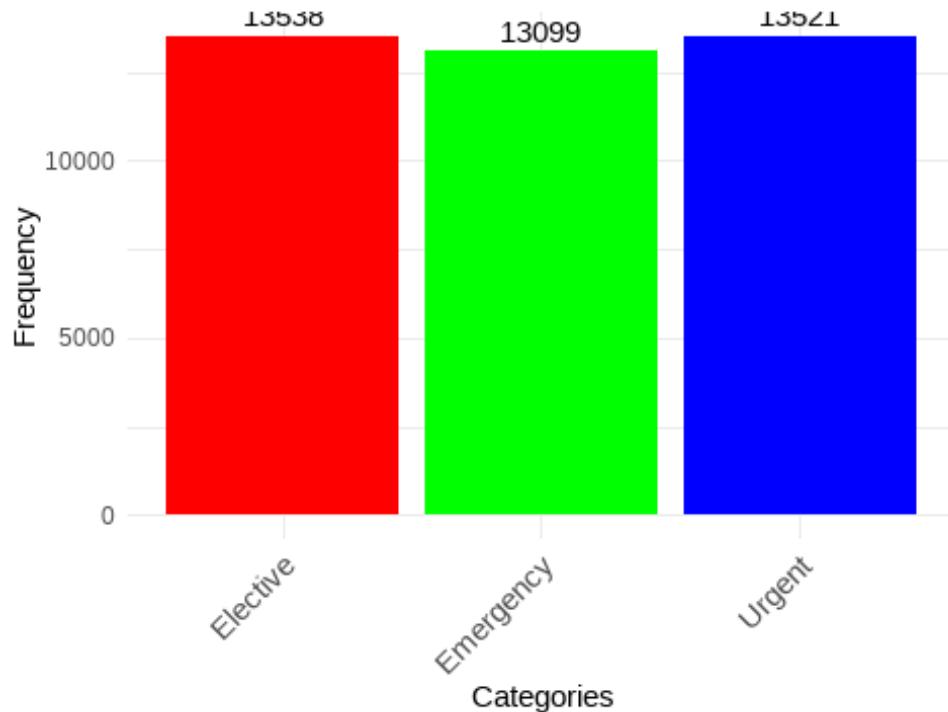
### Bar Chart - Insurance.Provider



```
## [1] "Bar chart is not applicable for Billing.Amount"  
## [1] "Bar chart is not applicable for Room.Number"
```

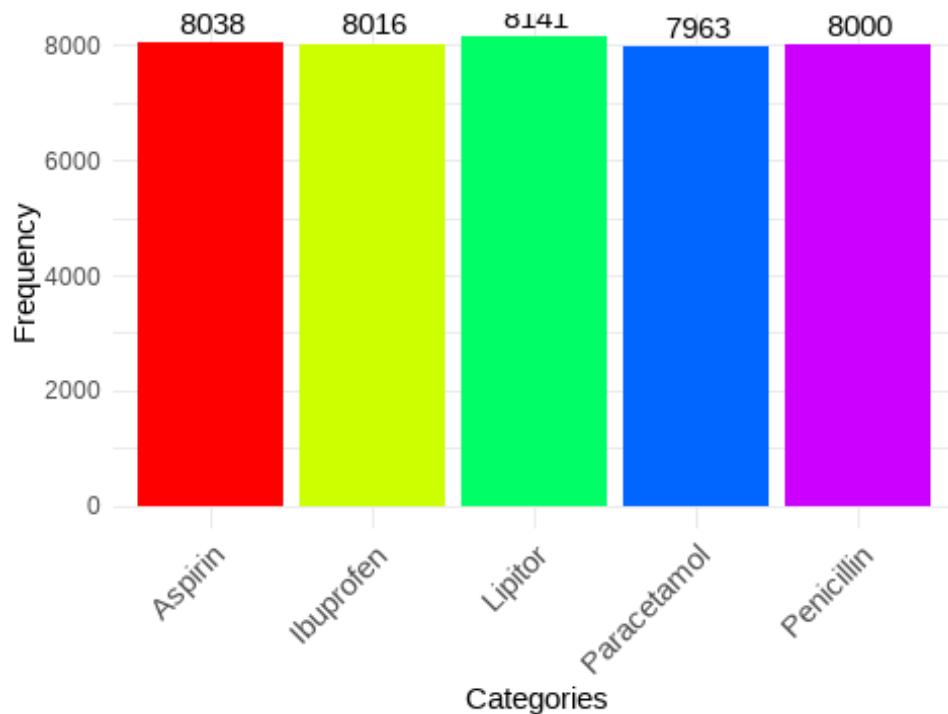
```
## Warning in `[<-.factor`(`*tmp*`, is.na(selected_var), value = "Unknown"):  
## invalid factor level, NA generated
```

**Bar Chart - Admission.Type**

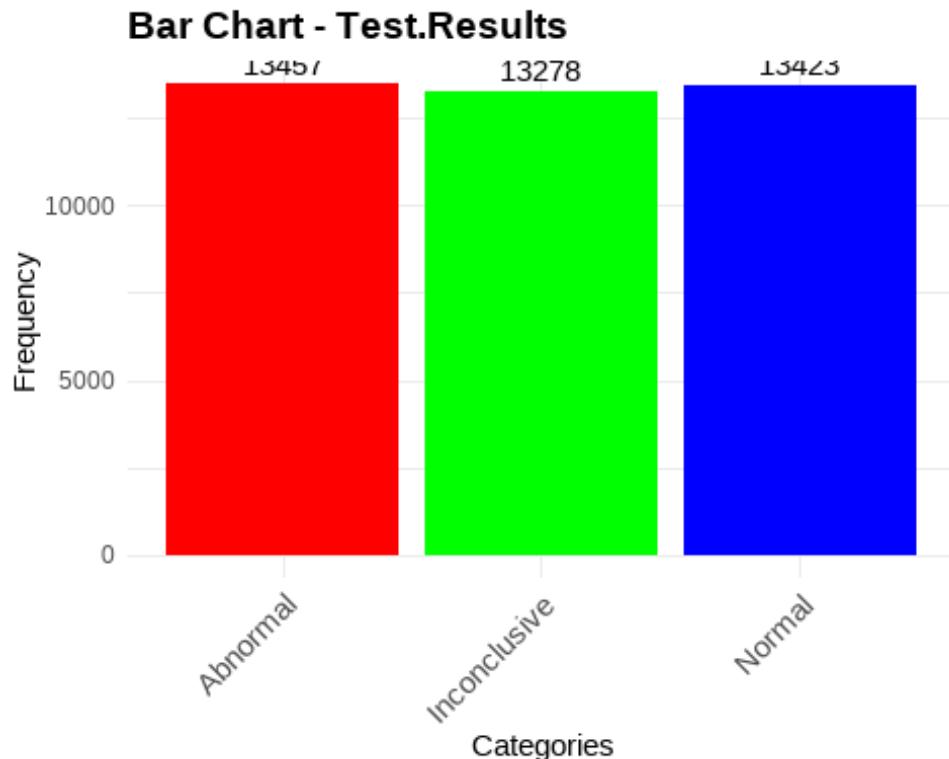


```
## Warning in `[<-.factor`(`*tmp*`, is.na(selected_var), value = "Unknown"):  
## invalid factor level, NA generated
```

**Bar Chart - Medication**



```
## Warning in `[<-.factor`(`*tmp*`, is.na(selected_var), value = "Unknown"):  
## invalid factor level, NA generated
```



Analysis: -

Verifies pie chart analysis as the output shows similar results

## Appendix 4: Variable Selection And Regression

As we can see from the EDA above the response variable (the outcome), admission type, is categorical data. Therefore, since the response variable only has 3 outcomes even though it has multiple independent variables (a factor that can cause changes in the response variable), so we are going to use multiclass logistic regression.

### Appendix 4.1: Loading Necessary Packages

```
library(MASS) # Used for advanced statistical modelling and variable selection  
  
##  
## Attaching package: 'MASS'  
  
## The following object is masked from 'package:dplyr':  
##  
##     select  
  
library(glmnet) # Used to fit the model  
  
## Loading required package: Matrix  
  
##  
## Attaching package: 'Matrix'
```

```

## The following objects are masked from 'package:tidy়':
##
##     expand, pack, unpack

## Loaded glmnet 4.1-10

library(Matrix) # For building Matrices
library(nnet)    # Used to fit the model
library(caret)   # Used for cross-validation, classification, building, tuning, and evaluating the model

## Loading required package: lattice

##
## Attaching package: 'caret'

## The following object is masked from 'package:purrr':
##
##     lift

library(hnp)      # Used to test for goodness of fit after variable selection

```

## Appendix 4.2: Variable Selection

### Appendix 4.2.1: Designing The Data For Model Fitting

```

# First, we make the predictors into a matrix and outcome as the factor outcome for admission type
x <- model.matrix(Admission.Type ~ ., data = health)[, -1]
y <- health$Admission.Type

```

### Appendix 4.2.2: Select Coefficients

```

## Pure LASSO
fit_lasso <- cv.glmnet(
  x, y,
  family = "multinomial",
  alpha = 1 # LASSO
)

# Best shrinkage parameter
fit_lasso$lambda.min

## [1] 0.006732435

# Selected coefficients
coef(fit_lasso, s = "lambda.min")

## $Elective
## 27 x 1 sparse Matrix of class "dgCMatrix"
##                                         lambda.min
## (Intercept)                      0.01140706
## Age                           .
## GenderMale                     .
## Blood.TypeA+                   .
## Blood.TypeAB-                  .
## Blood.TypeAB+                  .

```

```

## Blood.TypeB-
## Blood.TypeB+
## Blood.TypeO-
## Blood.TypeO+
## Medical.ConditionAsthma .
## Medical.ConditionCancer .
## Medical.ConditionDiabetes .
## Medical.ConditionHypertension .
## Medical.ConditionObesity .
## Insurance.ProviderBlue Cross .
## Insurance.ProviderCigna .
## Insurance.ProviderMedicare .
## Insurance.ProviderUnitedHealthcare .
## Billing.Amount .
## Room.Number .
## MedicationIbuprofen .
## MedicationLipitor .
## MedicationParacetamol .
## MedicationPenicillin .
## Test.ResultsInconclusive .
## Test.ResultsNormal .

##
## $Emergency
## 27 x 1 sparse Matrix of class "dgCMatrix"
##                               lambda.min
## (Intercept)                 -0.0215576
## Age .
## GenderMale .
## Blood.TypeA+ .
## Blood.TypeAB-
## Blood.TypeAB+ .
## Blood.TypeB- .
## Blood.TypeB+ .
## Blood.TypeO- .
## Blood.TypeO+ .
## Medical.ConditionAsthma .
## Medical.ConditionCancer .
## Medical.ConditionDiabetes .
## Medical.ConditionHypertension .
## Medical.ConditionObesity .
## Insurance.ProviderBlue Cross .
## Insurance.ProviderCigna .
## Insurance.ProviderMedicare .
## Insurance.ProviderUnitedHealthcare .
## Billing.Amount .
## Room.Number .
## MedicationIbuprofen .
## MedicationLipitor .
## MedicationParacetamol .
## MedicationPenicillin .
## Test.ResultsInconclusive .
## Test.ResultsNormal .

##
## $Urgent

```

```

## 27 x 1 sparse Matrix of class "dgCMatrix"
##                                     lambda.min
## (Intercept)                  0.01015054
## Age                         .
## GenderMale                   .
## Blood.TypeA+                 .
## Blood.TypeAB-                .
## Blood.TypeAB+                .
## Blood.TypeB-                 .
## Blood.TypeB+                 .
## Blood.TypeO-                 .
## Blood.TypeO+                 .
## Medical.ConditionAsthma     .
## Medical.ConditionCancer     .
## Medical.ConditionDiabetes   .
## Medical.ConditionHypertension .
## Medical.ConditionObesity    .
## Insurance.ProviderBlue Cross .
## Insurance.ProviderCigna     .
## Insurance.ProviderMedicare  .
## Insurance.ProviderUnitedHealthcare .
## Billing.Amount               .
## Room.Number                  .
## MedicationIbuprofen         .
## MedicationLipitor            .
## MedicationParacetamol        .
## MedicationPenicillin         .
## Test.ResultsInconclusive    .
## Test.ResultsNormal          .

```

*# Given that there are only values in the intercept, none of the variables were selected in this case*

*# But, since LASSO is quite an aggressive form of variable selection, we can try to use more ridge by reducing the alpha*

```

## A Balanced Elastic Net
fit_lasso <- cv.glmnet(
  x, y,
  family = "multinomial",
  alpha = 0.5
)

# Best shrinkage parameter
fit_lasso$lambda.min

## [1] 0.01346487

# Selected coefficients
coef(fit_lasso, s = "lambda.min")

## $Elective
## 27 x 1 sparse Matrix of class "dgCMatrix"
##                                     lambda.min
## (Intercept)                  0.01140706
## Age                         .

```

```

## GenderMale          .
## Blood.TypeA+        .
## Blood.TypeAB-       .
## Blood.TypeAB+       .
## Blood.TypeB-        .
## Blood.TypeB+        .
## Blood.TypeO-        .
## Blood.TypeO+        .
## Medical.ConditionAsthma   .
## Medical.ConditionCancer   .
## Medical.ConditionDiabetes  .
## Medical.ConditionHypertension .
## Medical.ConditionObesity   .
## Insurance.ProviderBlue Cross .
## Insurance.ProviderCigna    .
## Insurance.ProviderMedicare  .
## Insurance.ProviderUnitedHealthcare .
## Billing.Amount         .
## Room.Number            .
## MedicationIbuprofen     .
## MedicationLipitor       .
## MedicationParacetamol    .
## MedicationPenicillin     .
## Test.ResultsInconclusive .
## Test.ResultsNormal      .
##
## $Emergency
## 27 x 1 sparse Matrix of class "dgCMatrix"
##                                         lambda.min
## (Intercept)                         -0.0215576
## Age                                  .
## GenderMale                          .
## Blood.TypeA+                        .
## Blood.TypeAB-                       .
## Blood.TypeAB+                       .
## Blood.TypeB-                        .
## Blood.TypeB+                        .
## Blood.TypeO-                        .
## Blood.TypeO+                        .
## Medical.ConditionAsthma             .
## Medical.ConditionCancer            .
## Medical.ConditionDiabetes          .
## Medical.ConditionHypertension     .
## Medical.ConditionObesity          .
## Insurance.ProviderBlue Cross      .
## Insurance.ProviderCigna           .
## Insurance.ProviderMedicare        .
## Insurance.ProviderUnitedHealthcare .
## Billing.Amount                     .
## Room.Number                       .
## MedicationIbuprofen               .
## MedicationLipitor                 .
## MedicationParacetamol              .
## MedicationPenicillin               .

```

```

## Test.ResultsInconclusive      .
## Test.ResultsNormal          .
##
## $Urgent
## 27 x 1 sparse Matrix of class "dgCMatrix"
##                                         lambda.min
## (Intercept)                      0.01015054
## Age                                .
## GenderMale                          .
## Blood.TypeA+                        .
## Blood.TypeAB-                       .
## Blood.TypeAB+                       .
## Blood.TypeB-                        .
## Blood.TypeB+                        .
## Blood.TypeO-                        .
## Blood.TypeO+                        .
## Medical.ConditionAsthma            .
## Medical.ConditionCancer           .
## Medical.ConditionDiabetes         .
## Medical.ConditionHypertension     .
## Medical.ConditionObesity          .
## Insurance.ProviderBlue_Cross      .
## Insurance.ProviderCigna           .
## Insurance.ProviderMedicare        .
## Insurance.ProviderUnitedHealthcare .
## Billing.Amount                     .
## Room.Number                         .
## MedicationIbuprofen               .
## MedicationLipitor                 .
## MedicationParacetamol             .
## MedicationPenicillin              .
## Test.ResultsInconclusive          .
## Test.ResultsNormal                .

## We can see that there are selected variables now, but they are weak. So
## , we will not consider them when doing the predictions
## However, we can still use it for further interpretation later on. Hence
## , we will get the names of the selected variables for later

# Get exact names of the variables selected
selected_vars <- unique(unlist(lapply(
  coef(fit_lasso, s = "lambda.min"),
  function(m) rownames(m)[as.numeric(m) != 0]
)))
selected_vars <- setdiff(selected_vars, "(Intercept)")
selected_vars

## character(0)

```

## Appendix 4.3: Prediction

### Appendix 4.3.1: Splitting The Data

```
set.seed(100) # For reproducibility
```

```

# Randomly select 80% of the row as the split
split <- sample(1:nrow(health), 0.8 * nrow(health))
train_data <- health[split, ] # Use the 80% selected for training
test_data <- health[-split, ] # Use the remaining for testing

# Build matrices for training data
x_train <- sparse.model.matrix(Admission.Type ~ ., data = train_data)[,-1]
y_train <- train_data$Admission.Type

# Build matrices for test data
x_test <- sparse.model.matrix(Admission.Type ~ ., data = test_data)[,-1]
y_test <- test_data$Admission.Type

```

#### *Appendix 4.3.2: Fit the LASSO On Training Data And Test Data*

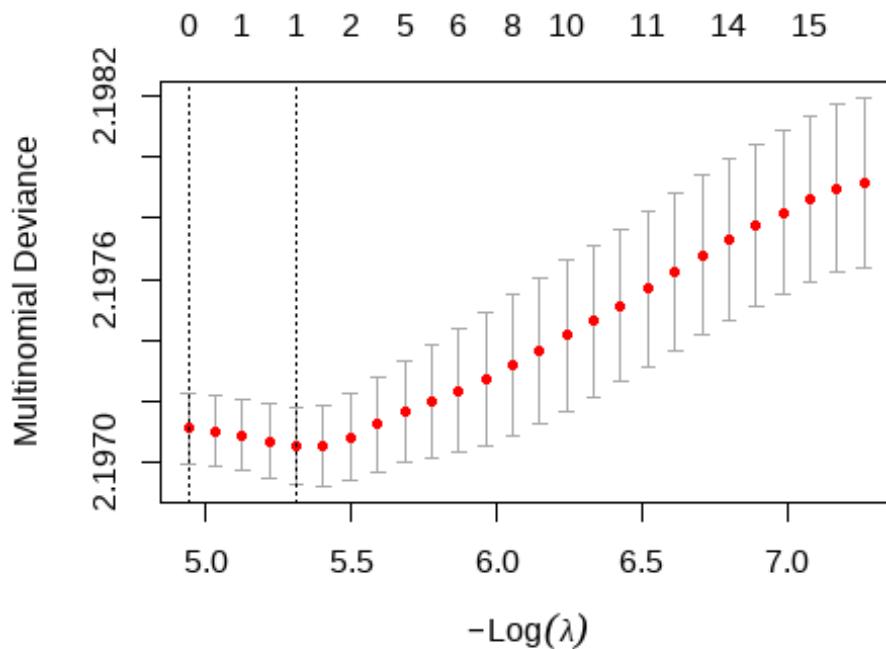
```

## Training Data
fit_lasso_train <- cv.glmnet(
  x_train, y_train,
  family = "multinomial",
  alpha = 1
)

## Test Data
fit_lasso_test <- cv.glmnet(
  x_test, y_test,
  family = "multi",
  alpha = 1
)

# Plotting cross-validation curve
plot(fit_lasso_train)

```



Analysis: -

Shows that the lambda.min (best fit model) and the lambda.1se (simpler model) will give similar outcomes. - So, we will just use the best fit model (which uses all the independent variables) going forward when we do predictions, seeing that there are not much benefits using a simpler model - However, we can still use the variable selection to see what other/further effects it might give to the model

#### *Appendix 4.3.3: Making Predictions Using The Full Model*

```
## Training data
pred_probs_train <- predict(fit_lasso_train, newx = x_train, s = "lambda.min",
                             type = "response")

## Test Data
pred_probs_test <- predict(fit_lasso_test, newx = x_test, s = "lambda.min",
                            type = "response")
```

#### *Appendix 4.3.4: Prediction Accuracy*

```
pred <- predict(fit_lasso_train, newx = x_test, s = "lambda.min", type = "class")
mean(pred == y_test)

## [1] 0.3367779
```

Analysis: Model is about 34% accurate, which is low. Though this was expected as all the variables were fairly evenly spread out, so we already expected that the predictions may not be very accurate

### *Appendix 4.4: Variable Selection Interpretation*

#### *Appendix 4.4.1: Summary of Variable Selection Fit*

```
# Refit the model with the selected variables
fit_refit <- multinom(
```

```

Admission.Type ~ Age +
  Blood.Type +
  Medical.Condition +
  Gender +
  Insurance.Provider +
  Medication,
  data = health
)

## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44088.289404
## iter 20 value 44087.563627
## iter 30 value 44085.872675
## iter 40 value 44084.599046
## final value 44084.511396
## converged

summary(fit_refit) # To see key outcomes of the model

## Call:
## multinom(formula = Admission.Type ~ Age + Blood.Type + Medical.Condition +
##           Gender + Insurance.Provider + Medication, data = health)
##
## Coefficients:
##             (Intercept)          Age          Blood.TypeA+          Blood.TypeAB-
## Emergency   -0.1188004  0.0010224649  0.003702696 -0.008871348   0.0310
## Urgent      -0.1681841  0.0008012514 -0.008373136 -0.037173542   0.0483
##             Blood.TypeB-          Blood.TypeB+          Blood.Type0-          Blood.Type0+
## Emergency  -0.008126562  0.06930849  -0.05008149  -0.02489878
## Urgent      0.010135709  0.04945616  -0.03906435  -0.05885322
##             Medical.ConditionAsthma          Medical.ConditionCancer
## Emergency      0.0009387206            -0.032504809
## Urgent        0.0027392519            0.008587153
##             Medical.ConditionDiabetes          Medical.ConditionHypertension
## Emergency      -0.00606062            -0.05548912
## Urgent        0.06541303            -0.02156108
##             Medical.ConditionObesity          GenderMale          Insurance.ProviderBlue C
##             ross
## Emergency      0.07211601  0.004905793   0.00997
## Urgent        0.02947333  0.048218580   0.05507
##             Insurance.ProviderCigna          Insurance.ProviderMedicare
## Emergency      0.02483461            0.02725148
## Urgent        0.11904184            0.10303487
##             Insurance.ProviderUnitedHealthcare          MedicationIbuprofen
## Emergency      -0.0008272783            0.02125015
## Urgent        0.0527720631            0.03707748
##             MedicationLipitor          MedicationParacetamol          MedicationPenicillin

```

```

## Emergency      0.06540553      -0.001774407      0.01635660
## Urgent        0.07000940      -0.017796399      0.03519527
##
## Std. Errors:
##             (Intercept)      Age Blood.TypeA+ Blood.TypeAB- Blood.Typ
eAB+
## Emergency  0.06581059  0.0006264306   0.04907518   0.04904322   0.0490
5811
## Urgent     0.06551693  0.0006215423   0.04869136   0.04875890   0.0485
0023
##             Blood.TypeB- Blood.TypeB+ Blood.Type0- Blood.Type0+
## Emergency  0.04915948  0.04907149   0.04908890   0.04904615
## Urgent      0.04859188  0.04873598   0.04856297   0.04880259
##             Medical.ConditionAsthma Medical.ConditionCancer
## Emergency    0.04231693           0.04228561
## Urgent       0.04216518           0.04192693
##             Medical.ConditionDiabetes Medical.ConditionHypertension
## Emergency    0.04227462           0.04224628
## Urgent       0.04176560           0.04192420
##             Medical.ConditionObesity GenderMale Insurance.ProviderBlue Cr
oss
## Emergency      0.04224930  0.02452960           0.03863
408
## Urgent        0.04233767  0.02433821           0.03863
961
##             Insurance.ProviderCigna Insurance.ProviderMedicare
## Emergency    0.03877443           0.03871923
## Urgent       0.03854424           0.03857228
##             Insurance.ProviderUnitedHealthcare MedicationIbuprofen
## Emergency      0.03878315           0.03877667
## Urgent        0.03874204           0.03844206
##             MedicationLipitor MedicationParacetamol MedicationPenicillin
## Emergency    0.03866484           0.03870538           0.03879883
## Urgent       0.03838771           0.03852948           0.03845511
##
## Residual Deviance: 88169.02
## AIC: 88261.02

```

Analysis: - Age: Looks that older patient are slightly more likely to have Emergency or Urgent admissions - Blood types: Admission types do vary by blood groups - Medical condition: Asthma patients are more likely to have more emergency or Urgent admissions. Diabetic patients do have more urgent admissions, though less for emergency admissions - Gender: Males shows to have slightly more likely urgent admissions compared to females - Insurance providers: Blue cross, cigna, medicare and United healthcare generally have higher number of emergency or urgent admitted patients with signa and medicare having stronger positive chance for urgent admissions - Medication: Lipitor shows some increasing effects on emergency and urgent admissions with Ibuprofen and Penicillin also having slightly lesser increasing effect on emergency and urgent admissions, but Paracetamol shows slight decreasing effect on urgent admissions

## *Appendix 4.4.2: Odds Ratios And Standard Errors*

To see the extent of the effects of the variables and how statistically significant are the variables

```
# Using the fit_refit we got just now, we get the log-odds coefficients to  
# find the odds ratios  
# Log-odds ratio  
coefs <- coef(fit_refit)  
  
# Compute odds ratios  
O.R. <- exp(coefs)  
  
# Standard errors  
s.e.s <- summary(fit_refit)$standard.errors
```

### *Appendix 4.4.3: Confidence Intervals (95%)*

```
lower_CI <- exp(coefs - 1.96 * s.e.s)  
upper_CI <- exp(coefs + 1.96 * s.e.s)
```

#### *Appendix 4.4.4: Combine Into A Table*

```

## 8 Emergency Blood.Type0-          -0.0501  0.951 0.0488  0.864
## 1.05
## 9 Emergency Blood.Type0+         -0.0249  0.975 0.0491  0.886
## 1.07
## 10 Emergency Medical.ConditionAsthma 0.000939 1.00  0.0485  0.910
## 1.10
## # i 36 more rows

```

Analysis: - Most predictors (independent variables) shows small or little effects on admission types - Blood type generally has weak and inconsistent effects - Insurance provider, being male, obese, diabetic does show some mild positive effect on emergency or urgent admissions compared to elective admissions.

#### *Appendix 4.4.5: Goodness of Fit*

```

# Make custom pearson residuals function
pearson_resid <- function(fit) {
  p <- fitted(fit) # Predicted probabilities
  y <- model.matrix(~ Admission.Type - 1, data = fit$model)

  res <- (y - p) / sqrt(p * (1 - p)) # Formula for residuals
  apply(res, 1, max) # Apply the maximum residual per observation
}

# Run HNP using the custom pearson residuals function and the refitted model with the selected variables
hnp(fit_refit, resid.fun = pearson_resid, nsim = 99)

## Multinomial model
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44075.340925
## iter 20 value 44071.753372
## iter 30 value 44069.219867
## final value 44068.721073
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44076.571002
## iter 20 value 44075.037201
## iter 30 value 44071.160380
## iter 40 value 44069.996397
## final value 44069.863806
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44087.005894
## iter 20 value 44084.315668
## iter 30 value 44083.750573
## iter 40 value 44083.364132
## final value 44083.225907
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44076.627809

```

```
## iter 20 value 44075.264697
## iter 30 value 44072.009659
## iter 40 value 44071.643064
## final value 44071.603774
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44067.014931
## iter 20 value 44065.540644
## iter 30 value 44061.766137
## iter 40 value 44060.066780
## final value 44059.617084
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44069.632705
## iter 20 value 44065.596578
## iter 30 value 44060.043300
## final value 44057.837188
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44072.437779
## iter 20 value 44071.668027
## iter 30 value 44069.622809
## final value 44068.073847
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44068.932884
## iter 20 value 44065.807030
## iter 30 value 44064.623776
## iter 40 value 44061.921070
## final value 44061.813370
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44069.417689
## iter 20 value 44067.875016
## iter 30 value 44062.432504
## iter 40 value 44060.871238
## final value 44060.471964
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44070.667653
## iter 20 value 44069.043138
## iter 30 value 44062.521381
## iter 30 value 44062.521163
## final value 44062.519937
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44063.419610
```

```
## iter 20 value 44061.431162
## iter 30 value 44056.410270
## iter 40 value 44055.254126
## final value 44055.160733
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44071.577922
## iter 20 value 44070.595757
## iter 30 value 44067.445209
## iter 40 value 44065.943522
## final value 44065.828709
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44067.686041
## iter 20 value 44066.337577
## iter 30 value 44062.504904
## iter 40 value 44060.894509
## final value 44060.356293
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44077.512266
## iter 20 value 44075.462154
## iter 30 value 44072.821963
## iter 40 value 44072.370502
## final value 44072.306131
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44071.774904
## iter 20 value 44070.349910
## iter 30 value 44067.294599
## iter 40 value 44066.223775
## final value 44065.530569
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44056.555422
## iter 20 value 44054.848188
## iter 30 value 44049.437300
## iter 40 value 44047.949280
## final value 44047.096203
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44075.631537
## iter 20 value 44071.973453
## iter 30 value 44069.538567
## iter 40 value 44068.635311
## final value 44068.483739
## converged
## # weights: 72 (46 variable)
```

```
## initial value 44118.072288
## iter 10 value 44068.757904
## iter 20 value 44067.546772
## iter 30 value 44062.902570
## iter 40 value 44061.911919
## final value 44061.758529
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44060.171701
## iter 20 value 44056.514251
## iter 30 value 44047.876877
## iter 40 value 44046.604448
## final value 44046.521418
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44061.823757
## iter 20 value 44056.311417
## iter 30 value 44052.477486
## iter 40 value 44051.773640
## final value 44051.662848
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44067.389339
## iter 20 value 44065.593918
## iter 30 value 44059.626633
## iter 40 value 44058.859107
## final value 44058.468365
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44073.343424
## iter 20 value 44070.385618
## iter 30 value 44065.764302
## iter 40 value 44064.594415
## final value 44064.455951
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44068.535499
## iter 20 value 44066.759198
## iter 30 value 44062.070993
## iter 40 value 44059.780700
## final value 44059.422211
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44073.234343
## iter 20 value 44068.379968
## iter 30 value 44063.428420
## final value 44062.922266
## converged
```

```
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44072.864321
## iter 20 value 44069.991529
## iter 30 value 44062.797188
## final value 44061.504653
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44066.512369
## iter 20 value 44064.614015
## iter 30 value 44060.042403
## iter 40 value 44057.064353
## final value 44056.667725
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44049.712560
## iter 20 value 44048.271933
## iter 30 value 44042.681786
## iter 40 value 44041.201765
## final value 44041.120788
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44066.575687
## iter 20 value 44065.541726
## iter 30 value 44062.681792
## iter 40 value 44061.790749
## final value 44061.734413
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44066.631688
## iter 20 value 44065.078746
## iter 30 value 44060.359896
## iter 40 value 44057.107185
## final value 44056.809745
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44082.080254
## iter 20 value 44080.589386
## iter 30 value 44076.761391
## iter 40 value 44076.242485
## final value 44076.182507
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44084.000178
## iter 20 value 44081.872036
## iter 30 value 44078.574136
## iter 40 value 44077.440425
## final value 44077.334000
```

```
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44078.548816
## iter 20 value 44077.277790
## iter 30 value 44074.830396
## iter 40 value 44073.233230
## final value 44073.166428
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44066.991060
## iter 20 value 44063.998185
## iter 30 value 44058.858721
## iter 40 value 44057.891634
## final value 44057.852031
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44071.181098
## iter 20 value 44070.228483
## iter 30 value 44066.580838
## iter 40 value 44065.554677
## final value 44065.445861
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44070.758068
## iter 20 value 44067.443849
## iter 30 value 44059.625232
## iter 40 value 44057.646969
## final value 44057.303927
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44074.001523
## iter 20 value 44070.446099
## iter 30 value 44067.726160
## iter 40 value 44066.365880
## final value 44066.238586
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44065.054266
## iter 20 value 44063.472424
## iter 30 value 44060.305504
## iter 40 value 44059.216010
## final value 44059.034406
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44079.358943
## iter 20 value 44077.467468
## iter 30 value 44069.921393
```

```
## iter 40 value 44068.638661
## final value 44068.250659
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44079.657018
## iter 20 value 44077.681879
## iter 30 value 44071.770473
## iter 40 value 44070.904441
## final value 44070.288746
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44068.163989
## iter 20 value 44067.098990
## iter 30 value 44065.016576
## final value 44064.317226
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44081.056844
## iter 20 value 44077.734813
## iter 30 value 44074.863752
## iter 40 value 44074.417476
## final value 44074.286743
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44080.418453
## iter 20 value 44078.961204
## iter 30 value 44075.663620
## final value 44073.561916
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44085.315094
## iter 20 value 44083.776273
## iter 30 value 44079.653592
## iter 40 value 44078.225125
## final value 44078.087521
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44070.815900
## iter 20 value 44069.550461
## iter 30 value 44065.566520
## iter 40 value 44064.948042
## final value 44064.711083
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44064.550482
## iter 20 value 44062.236585
## iter 30 value 44055.749735
```

```
## iter 40 value 44054.648465
## final value 44054.265013
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44078.215236
## iter 20 value 44076.817989
## iter 30 value 44072.781973
## iter 40 value 44070.582814
## final value 44070.442672
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44058.896265
## iter 20 value 44055.512049
## iter 30 value 44051.541606
## iter 40 value 44047.215682
## iter 50 value 44046.526523
## iter 50 value 44046.526282
## iter 50 value 44046.526280
## final value 44046.526280
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44050.444315
## iter 20 value 44047.339199
## iter 30 value 44043.954353
## iter 40 value 44043.128538
## final value 44042.963859
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44072.842578
## iter 20 value 44071.007199
## iter 30 value 44067.863960
## final value 44067.263116
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44076.689179
## iter 20 value 44075.535957
## iter 30 value 44072.671424
## iter 40 value 44071.080655
## final value 44070.396830
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44061.278181
## iter 20 value 44057.911302
## iter 30 value 44054.031211
## final value 44052.596024
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
```

```
## iter 10 value 44078.898473
## iter 20 value 44075.800749
## iter 30 value 44073.973093
## iter 40 value 44073.576875
## iter 40 value 44073.576710
## iter 40 value 44073.576708
## final value 44073.576708
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44069.076976
## iter 20 value 44066.726965
## iter 30 value 44058.813616
## final value 44057.411685
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44072.656339
## iter 20 value 44068.848647
## iter 30 value 44064.975411
## iter 40 value 44063.237496
## final value 44063.122627
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44064.964896
## iter 20 value 44063.213598
## iter 30 value 44059.343359
## iter 30 value 44059.343269
## final value 44058.629583
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44061.733613
## iter 20 value 44059.657841
## iter 30 value 44052.917954
## final value 44051.003443
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44068.171663
## iter 20 value 44064.473629
## iter 30 value 44060.519440
## iter 40 value 44059.160630
## final value 44059.107793
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44075.155221
## iter 20 value 44074.298026
## iter 30 value 44071.711654
## final value 44071.023013
## converged
## # weights: 72 (46 variable)
```

```
## initial value 44118.072288
## iter 10 value 44057.384515
## iter 20 value 44055.481782
## iter 30 value 44051.080136
## iter 40 value 44049.559464
## final value 44049.339890
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44061.082043
## iter 20 value 44060.144349
## iter 30 value 44057.463552
## final value 44056.670878
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44077.146362
## iter 20 value 44076.051422
## iter 30 value 44072.197089
## iter 40 value 44071.417172
## final value 44071.118751
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44070.921554
## iter 20 value 44070.049918
## iter 30 value 44067.019365
## iter 40 value 44065.359104
## final value 44065.196246
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44069.542899
## iter 20 value 44067.324847
## iter 30 value 44062.354447
## iter 40 value 44061.546456
## final value 44061.395451
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44077.590018
## iter 20 value 44074.201367
## iter 30 value 44071.151460
## iter 40 value 44070.378815
## final value 44070.225482
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44068.519232
## iter 20 value 44066.341881
## iter 30 value 44063.971351
## iter 40 value 44062.990519
## final value 44062.849420
## converged
```

```

## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44067.743396
## iter 20 value 44065.921025
## iter 30 value 44060.427568
## iter 40 value 44059.474461
## final value 44059.165831
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44077.504031
## iter 20 value 44076.200983
## iter 30 value 44074.824056
## iter 40 value 44074.657879
## final value 44074.470199
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44075.038364
## iter 20 value 44072.514498
## iter 30 value 44070.684630
## iter 40 value 44069.485823
## final value 44069.417410
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44063.963832
## iter 20 value 44062.723061
## iter 30 value 44057.920657
## iter 40 value 44057.211343
## final value 44056.944763
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44053.698289
## iter 20 value 44050.184289
## iter 30 value 44040.938256
## iter 40 value 44039.863165
## final value 44039.675985
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44060.483880
## iter 20 value 44057.888836
## iter 30 value 44053.780136
## iter 40 value 44052.709883
## final value 44052.564184
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44067.321494
## iter 20 value 44063.493851
## iter 30 value 44062.182694
## final value 44061.654625

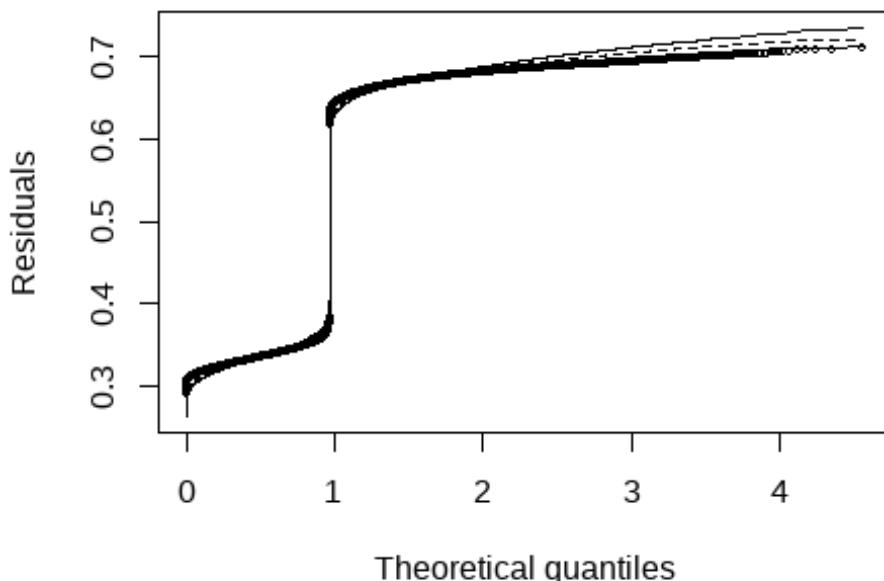
```

```
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44080.741294
## iter 20 value 44078.049540
## iter 30 value 44076.462595
## final value 44075.852037
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44065.141622
## iter 20 value 44062.977424
## iter 30 value 44055.339005
## final value 44054.807333
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44055.169798
## iter 20 value 44052.868116
## iter 30 value 44048.585627
## iter 40 value 44047.553983
## final value 44047.533450
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44065.435604
## iter 20 value 44060.925202
## iter 30 value 44050.133971
## iter 40 value 44047.483379
## final value 44047.321891
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44081.186634
## iter 20 value 44080.112917
## iter 30 value 44076.892506
## iter 40 value 44076.153427
## final value 44075.844043
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44072.705572
## iter 20 value 44069.222560
## iter 30 value 44064.468477
## iter 40 value 44062.936966
## final value 44062.805570
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44064.566755
## iter 20 value 44062.186559
## iter 30 value 44055.348274
## iter 30 value 44055.347885
## iter 40 value 44054.364095
```

```
## iter 50 value 44054.220968
## final value 44054.107057
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44056.780219
## iter 20 value 44054.761082
## iter 30 value 44046.837431
## final value 44044.712288
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44057.984806
## iter 20 value 44056.710655
## iter 30 value 44051.422328
## iter 40 value 44050.233596
## final value 44049.159781
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44064.366007
## iter 20 value 44062.799838
## iter 30 value 44057.935209
## final value 44056.457547
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44062.162730
## iter 20 value 44060.477520
## iter 30 value 44055.365679
## iter 40 value 44054.398827
## final value 44054.310018
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44081.816096
## iter 20 value 44078.645834
## iter 30 value 44075.412779
## iter 40 value 44073.874449
## final value 44073.517687
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44076.744257
## iter 20 value 44075.347639
## iter 30 value 44071.313075
## iter 30 value 44071.312776
## iter 40 value 44070.919943
## iter 50 value 44070.879240
## final value 44070.770461
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44061.225709
```

```
## iter 20 value 44056.976510
## iter 30 value 44054.345050
## iter 40 value 44053.289266
## final value 44053.200305
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44062.746619
## iter 20 value 44060.029300
## iter 30 value 44056.402843
## iter 40 value 44055.659915
## final value 44055.564488
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44066.040484
## iter 20 value 44064.542117
## iter 30 value 44060.886335
## iter 40 value 44058.786071
## final value 44058.505066
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44070.823482
## iter 20 value 44068.172809
## iter 30 value 44062.668863
## iter 40 value 44061.761779
## final value 44061.703977
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44070.197435
## iter 20 value 44064.602396
## iter 30 value 44059.162391
## iter 40 value 44057.781658
## final value 44057.568548
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44062.324105
## iter 20 value 44061.058670
## iter 30 value 44058.170871
## iter 40 value 44057.426514
## final value 44057.377398
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44085.808409
## iter 20 value 44084.961051
## iter 30 value 44082.082812
## iter 40 value 44081.725660
## final value 44081.673952
## converged
## # weights: 72 (46 variable)
```

```
## initial value 44118.072288
## iter 10 value 44053.087323
## iter 20 value 44051.936328
## iter 30 value 44047.972105
## final value 44046.453936
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44053.997616
## iter 20 value 44052.391576
## iter 30 value 44047.395972
## iter 40 value 44045.955236
## final value 44045.655580
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44059.149409
## iter 20 value 44058.234107
## iter 30 value 44056.250331
## final value 44054.230510
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44066.631834
## iter 20 value 44064.879630
## iter 30 value 44059.747550
## iter 40 value 44058.441388
## final value 44058.371771
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44077.325481
## iter 20 value 44075.965047
## iter 30 value 44072.199063
## iter 40 value 44071.127712
## final value 44071.004935
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44064.109153
## iter 20 value 44058.781300
## iter 30 value 44048.464344
## iter 40 value 44047.361413
## final value 44046.970481
## converged
## # weights: 72 (46 variable)
## initial value 44118.072288
## iter 10 value 44074.030371
## iter 20 value 44072.917034
## iter 30 value 44070.423234
## iter 40 value 44069.878783
## final value 44069.763830
## converged
```



Analysis: -

Given that this is fitted using a multinomial logistic regression model, the vertical line in the output is expected as the response variable is categorical and the residuals can only take a few values - Though with the smooth curve, it can be seen that the largest residuals are slightly outside the envelope which indicates there are missing predictors - This indicates that the fitted model after variable selection is a moderate-to-poor fit, which justifies the previous decision of using all the variables for the model instead of the model after variable selection

## Appendix 4.5: Cross-Validation

### Appendix 4.5.1: K-Fold Cross Validation

```
set.seed(100) # For reproducibility

## 5-fold cross-validation
# Define cross-validation settings
cv_control_5 <- trainControl(
  method = "cv",
  number = 5,
)

# Train multinomial Logistic regression with CV
model_cv5 <- train(
  Admission.Type ~.,
  data = health,
  method = "multinom",    # Tells that its a multinomial Logistic regression
n
  trControl = cv_control_5,
  trace = FALSE
)
```

```

# Print CV results
print(model_cv5)

## Penalized Multinomial Regression
##
## 40158 samples
##     9 predictor
##      3 classes: 'Elective', 'Emergency', 'Urgent'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold)
## Summary of sample sizes: 32126, 32127, 32127, 32127, 32125
## Resampling results across tuning parameters:
##
##     decay  Accuracy   Kappa
##     0e+00  0.3351761 -0.0003591819
##     1e-04  0.3351761 -0.0003595632
##     1e-01  0.3352259 -0.0002655195
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was decay = 0.1.

# Access resampling results, e.g., accuracy
model_cv5$resample$Accuracy

## [1] 0.3310508 0.3379405 0.3360727 0.3368605 0.3342050

mean(model_cv5$resample$Accuracy)

## [1] 0.3352259

## Next, we try 10-fold cross-validation to see if there are any improvements
set.seed(100) # For reproducibility

## 10-fold cross-validation
# Define cross-validation settings
cv_control_10 <- trainControl(
  method = "cv",
  number = 10,
)

# Train multinomial Logistic regression with CV
model_cv10 <- train(
  Admission.Type ~.,
  data = health,
  method = "multinom",    # Tells that its a multinomial Logistic regression
n
  trControl = cv_control_10,
  trace = FALSE
)

# Print CV results
print(model_cv10)

```

```

## Penalized Multinomial Regression
##
## 40158 samples
##      9 predictor
##      3 classes: 'Elective', 'Emergency', 'Urgent'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 36142, 36142, 36142, 36141, 36143, 36144, ...
## Resampling results across tuning parameters:
##
##   decay  Accuracy  Kappa
##   0e+00  0.3407289  0.007869052
##   1e-04  0.3407538  0.007905297
##   1e-01  0.3408534  0.008054431
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was decay = 0.1.

# Access resampling results, e.g., accuracy
model_cv10$resample$Accuracy

## [1] 0.3249502 0.3453685 0.3452825 0.3336653 0.3338316 0.3484433 0.3361
554
## [8] 0.3505976 0.3446215 0.3456175

mean(model_cv10$resample$Accuracy)

## [1] 0.3408534

```

Analysis 5-fold CV: - Average accuracy is about 33.5% which is very low - Kappa is negative, so the model is no better than pure guessing using 5-fold CV - Shows that model has very weak predictive power

Analysis 10-fold CV: - Average accuracy is about 33.8% which is only slightly better than 5-fold CV - Kappa is positive now, but (~0.004) still less than 0.05, so model is not much better than pure guessing - Still shows that model has very weak predictive power

#### *Appendix 4.5.2: Confusion Matrix*

```

confusionMatrix(factor(pred, levels = levels(y_test)), y_test)

## Confusion Matrix and Statistics
##
##             Reference
## Prediction  Elective Emergency Urgent
##   Elective      2237     2194    2302
##   Emergency      0        0       0
##   Urgent         427     404    468
##
## Overall Statistics
##
##               Accuracy : 0.3368
##                 95% CI : (0.3264, 0.3472)
##   No Information Rate : 0.3449

```

```

##      P-Value [Acc > NIR] : 0.9381
##
##          Kappa : 0.0045
##
##  Mcnemar's Test P-Value : <2e-16
##
## Statistics by Class:
##
##          Class: Elective Class: Emergency Class: Urgent
## Sensitivity           0.8397       0.0000    0.16895
## Specificity            0.1624       1.0000    0.84208
## Pos Pred Value         0.3322        NaN     0.36028
## Neg Pred Value         0.6713       0.6765    0.65810
## Prevalence              0.3317       0.3235    0.34487
## Detection Rate          0.2785       0.0000    0.05827
## Detection Prevalence    0.8383       0.0000    0.16173
## Balanced Accuracy        0.5011       0.5000    0.50551

```

Analysis: - About 33.7% is correct, which is what we have expected from the previous results - Kappa is 0.0045, which confirms the k-fold cross-validation - Though we can note here, that the model completely fails to predict emergency cases, but the model still performs at chance level for every class of admission type - Results indicate that the available predictors are not suitable or insufficient to distinguish admission types, particularly emergency admissions

## Appendix 5: Classification

### Appendix 5.1: Loading Necessary Packages

```

library(e1071)          # For data mining and classification

##
## Attaching package: 'e1071'

## The following objects are masked from 'package:dlookr':
##
##      kurtosis, skewness

## The following object is masked from 'package:ggplot2':
##
##      element

library(randomForest) # For random forests

## randomForest 4.7-1.2

## Type rfNews() to see new features/changes/bug fixes.

##
## Attaching package: 'randomForest'

## The following object is masked from 'package:dplyr':
##
##      combine

```

```

## The following object is masked from 'package:ggplot2':
##
##      margin

library(class)          # For KNN
library(kernlab)        # To back-end for SVM (Radial)

##
## Attaching package: 'kernlab'

## The following object is masked from 'package:scales':
##
##      alpha

## The following object is masked from 'package:purrr':
##
##      cross

## The following object is masked from 'package:ggplot2':
##
##      alpha

```

## Appendix 5.2: Resplit The Data

```

set.seed(100) # For reproducibility

# Randomly select 80% of the row as the split
split <- sample(1:nrow(health), 0.8 * nrow(health))
train_data <- health[split, ] # Use the 80% selected for training
test_data <- health[-split, ] # Use the remaining for testing

```

## Appendix 5.3: Reprocessing And Train Control

```

# Set the instructions on how to train and evaluate the model
train_control <- trainControl(
  method = "cv",
  number = 5,
  classProbs = TRUE,
  summaryFunction = multiClassSummary # Multiclass performance matrix
)

```

## Appendix 5.4: Random Forest

### Appendix 5.4.1: Random Forest Model

```

# Train the model
rf_model <- train(
  Admission.Type ~ .,
  data = train_data, # Use the training data set
  method = "ranger", # Fits an efficient RF, supports Large datasets and m
  ulticlass problems efficiently
  trControl = train_control, # Applies the previously set instructions
  tuneLength = 5 # Automatically tries different mtry values and selects t
  he best one
)

```

```

## Growing trees.. Progress: 66%. Estimated remaining time: 15 seconds.
## Growing trees.. Progress: 64%. Estimated remaining time: 17 seconds.
## Growing trees.. Progress: 95%. Estimated remaining time: 1 seconds.
## Growing trees.. Progress: 76%. Estimated remaining time: 10 seconds.
## Growing trees.. Progress: 81%. Estimated remaining time: 7 seconds.
## Growing trees.. Progress: 94%. Estimated remaining time: 2 seconds.
## Growing trees.. Progress: 79%. Estimated remaining time: 8 seconds.
## Growing trees.. Progress: 74%. Estimated remaining time: 11 seconds.
## Growing trees.. Progress: 82%. Estimated remaining time: 6 seconds.
## Growing trees.. Progress: 83%. Estimated remaining time: 6 seconds.
## Growing trees.. Progress: 79%. Estimated remaining time: 8 seconds.
## Growing trees.. Progress: 89%. Estimated remaining time: 3 seconds.
## Growing trees.. Progress: 84%. Estimated remaining time: 6 seconds.
## Growing trees.. Progress: 100%. Estimated remaining time: 0 seconds.
## Growing trees.. Progress: 73%. Estimated remaining time: 11 seconds.

```

#### *Appendix 5.4.2: Predictions and Evaluation*

```

rf_pred <- predict(rf_model, newdata = test_data) # Predict
rf_acc <- mean(rf_pred == test_data$Admission.Type) # Accuracy
rf_cm <- confusionMatrix(rf_pred, test_data$Admission.Type) # Overall mode
l evaluation

print("Random Forest Accuracy:")
## [1] "Random Forest Accuracy:"
```

```

print(rf_acc)
## [1] 0.3306773

print("Random Forest Confusion Matrix:")
## [1] "Random Forest Confusion Matrix:"
```

```

print(rf_cm)

## Confusion Matrix and Statistics
##
##          Reference
## Prediction Elective Emergency Urgent
##   Elective      1201       1237     1276
##   Emergency     503        484      523
##   Urgent         960        877      971
##
## Overall Statistics
##
##                 Accuracy : 0.3307
##                           95% CI : (0.3204, 0.3411)
##   No Information Rate : 0.3449
##   P-Value [Acc > NIR] : 0.9965
##
##                 Kappa : -0.0061
##
## McNemar's Test P-Value : <2e-16
##
```

```

## Statistics by Class:
##
##          Class: Elective Class: Emergency Class: Urgent
## Sensitivity           0.4508      0.18630    0.3505
## Specificity           0.5319      0.81119    0.6509
## Pos Pred Value        0.3234      0.32053    0.3458
## Neg Pred Value        0.6612      0.67587    0.6556
## Prevalence            0.3317      0.32346    0.3449
## Detection Rate        0.1495      0.06026    0.1209
## Detection Prevalence  0.4624      0.18800    0.3496
## Balanced Accuracy     0.4913      0.49874    0.5007

```

Analysis: - Accuracy is still about 33.1%, which is almost the same as the regression model, no improvements - Though random forest does predict all 3 classes of admission types (unlike the multinomial model), but the predictions are almost uniformly spread - Kappa is negative here, which indicates that its worse than random guessing an ocnfirms that the predictions are essentially random

## Appendix 5.5: K-Nearest Neighbors (KNN)

### Appendix 5.5.1: KNN Model

```

knn_model <- train(
  Admission.Type ~ .,
  data = train_data, # Use training data
  method = "knn", # Use K-Nearest Neighbors classification
  preProcess = c("center", "scale"), # Standardise predictors
  trControl = train_control,# Applies the previously set instructions
  tuneLength = 10 # Automatically tries multiple k values and select the best one
)

```

### Appendix 5.5.2: Predictions and Evaluation

```

knn_pred <- predict(knn_model, newdata = test_data) # Predict
knn_acc <- mean(knn_pred == test_data$Admission.Type) # Accuracy
knn_cm <- confusionMatrix(knn_pred, test_data$Admission.Type) # Overall model evaluation

print("k-NN Accuracy:")
## [1] "k-NN Accuracy:"
print(knn_acc)
## [1] 0.3290588

print("k-NN Confusion Matrix:")
## [1] "k-NN Confusion Matrix:"
print(knn_cm)

## Confusion Matrix and Statistics
##
##          Reference
## Prediction   Elective Emergency Urgent

```

```

##   Elective      938      927      993
##   Emergency    839      809      881
##   Urgent        887      862      896
##
## Overall Statistics
##
##               Accuracy : 0.3291
##                   95% CI : (0.3188, 0.3395)
##   No Information Rate : 0.3449
##   P-Value [Acc > NIR] : 0.9987
##
##               Kappa : -0.0066
##
## McNemar's Test P-Value : 0.0143
##
## Statistics by Class:
##
##                               Class: Elective Class: Emergency Class: Urgent
## Sensitivity                  0.3521      0.3114      0.3235
## Specificity                  0.6423      0.6835      0.6676
## Pos Pred Value                0.3282      0.3199      0.3388
## Neg Pred Value                0.6664      0.6749      0.6521
## Prevalence                    0.3317      0.3235      0.3449
## Detection Rate                 0.1168      0.1007      0.1116
## Detection Prevalence          0.3558      0.3149      0.3293
## Balanced Accuracy              0.4972      0.4974      0.4955

```

Analysis: - Accuracy is about 32.9%, which is no better than random forests - Again, most predictions are almost evenly spread across all admission types - Again, kappa is negative, so model is no better than random guessing

## Appendix 5.6: Linear Discriminant Analysis (LDA)

### Appendix 5.6.1: LDA Model

```

lda_model <- train(
  Admission.Type ~ .,
  data = train_data, # Use training data
  method = "lda", # Use LDA classification
  trControl = train_control # Applies the previously set instructions
)

```

### Appendix 5.6.4: Predictions And Evaluation

```

lda_pred <- predict(lda_model, newdata = test_data) # Predict
lda_acc <- mean(lda_pred == test_data$Admission.Type) # Accuracy
lda_cm <- confusionMatrix(lda_pred, test_data$Admission.Type) # Overall model evaluation

print("LDA Accuracy:")
## [1] "LDA Accuracy:"
print(lda_acc)
## [1] 0.3359064

```

```

print("LDA Confusion Matrix:")
## [1] "LDA Confusion Matrix:"
```

```

print(lda_cm)
```

```

## Confusion Matrix and Statistics
##
##             Reference
## Prediction   Elective Emergency Urgent
##   Elective      1298       1284     1343
##   Emergency     371        335      362
##   Urgent         995       979     1065
##
## Overall Statistics
##
##                 Accuracy : 0.3359
##                 95% CI : (0.3256, 0.3464)
##   No Information Rate : 0.3449
##   P-Value [Acc > NIR] : 0.9558
##
##                 Kappa : 5e-04
##
## McNemar's Test P-Value : <2e-16
##
## Statistics by Class:
##
##                                Class: Elective Class: Emergency Class: Urgent
## Sensitivity                  0.4872      0.12895      0.3845
## Specificity                  0.5106      0.86511      0.6249
## Pos Pred Value                0.3307      0.31367      0.3504
## Neg Pred Value                0.6674      0.67504      0.6585
## Prevalence                    0.3317      0.32346      0.3449
## Detection Rate                 0.1616      0.04171      0.1326
## Detection Prevalence          0.4887      0.13297      0.3784
## Balanced Accuracy              0.4989      0.49703      0.5047

```

Analysis: - Accuracy is about 33.6%, similar to previous results - Again, most predictions are almost evenly spread across all admission types - Kappa is positive here, but (0.0003) still less than 0.05, effectively no better than random guessing

- Overall, every model gives a chance performance, so the problem is not the model choice.
- It may be that the predictors do not contain much usable information to distinguish admission types

## Appendix 6: Deep Learning

### Appendix 6.1: Loading Necessary Packages

```

library(keras3) # Used for deep Learning
library(tensorflow) # Allows to run custom deep-Learning computations
```

```

##
## Attaching package: 'tensorflow'
```

```

## The following objects are masked from 'package:keras3':
##
##     set_random_seed, shape

## The following object is masked from 'package:caret':
##
##     train

```

## Appendix 6.1: Resplit The Data

```

set.seed(100)# For reproducibility

# Randomly select 80% of the row as the split
split <- sample(1:nrow(health), 0.8 * nrow(health))
train_data <- health[split, ] # Use the 80% selected for training
test_data <- health[-split, ] # Use the remaining for testing

```

## Appendix 6.2: Reprocessing Data

```

# Separate predictors and response variables
# Using training data
x_train <- model.matrix(Admission.Type ~ ., data = train_data)[, -1] # Remove the intercept

# Using test data
x_test <- model.matrix(Admission.Type ~ ., data = test_data)[, -1] # Remove the intercept

```

## Appendix 6.3: Scale Predictors

```

# Standardising independent variables x
x_train <- scale(x_train)
x_test <- scale(
  x_test,
  center = attr(x_train, "scaled:center"), # Mean of each column
  scale = attr(x_train, "scaled:scale") # Standard deviation of each column
)
# Defining admission type as the response variable y
y_train <- train_data$Admission.Type
y_test <- test_data$Admission.Type

# Converting training labels to categorical
y_train_cat <- tf$keras$utils$to_categorical(as.integer(y_train) - 1) # -1 to ensure classes starts at 0

# Converting test labels to categorical
y_test_cat <- tf$keras$utils$to_categorical(as.integer(y_test) - 1) # -1 to ensure classes starts at 0

num_classes <- ncol(y_train_cat) # Gives the number of unique classes

```

## Appendix 6.4: Define Neural Network

```
model <- keras_model_sequential() %>%
  # Add first layer
  layer_dense(units = 32, activation = "relu") %>%
  
  # Dropout after the first Layer to prevent overfitting (drop 30% of neurons)
  layer_dropout(rate = 0.3) %>%
  
  # Add second layer
  layer_dense(units = 16, activation = "relu") %>%
  
  # Dropout after second Layer to prevent overfitting in deeper Layers (drop 30% of neurons)
  layer_dropout(rate = 0.3) %>%
  
  # Output Layer (final classification layer)
  layer_dense(units = num_classes, activation = "softmax")
```

## Appendix 6.7: Compile And Train Model

```
# Compile model
model %>% compile(
  optimizer = optimizer_adam(learning_rate = 0.001), # Use Adam optimiser to adapt Learning rates automatically for faster and stable training
  loss = "categorical_crossentropy", # Measures the difference between predicted probabilities and true labels
  metrics = "accuracy" # Tracks accuracy
)

# Train model
history <- model %>% fit(
  x_train, y_train_cat,
  validation_split = 0.2, # Reserve 20% of training data for validation to monitor overfitting
  epochs = 50, # Number of full passes through the training data
  batch_size = 32 # Number of samples processed before updating the weight
)
## Epoch 1/50
## 804/804 - 2s - 3ms/step - accuracy: 0.3332 - loss: 1.1181 - val_accuracy: 0.3340 - val_loss: 1.0986
## Epoch 2/50
## 804/804 - 1s - 1ms/step - accuracy: 0.3460 - loss: 1.0994 - val_accuracy: 0.3372 - val_loss: 1.0985
## Epoch 3/50
## 804/804 - 1s - 1ms/step - accuracy: 0.3362 - loss: 1.0992 - val_accuracy: 0.3368 - val_loss: 1.0985
## Epoch 4/50
## 804/804 - 1s - 1ms/step - accuracy: 0.3416 - loss: 1.0986 - val_accuracy: 0.3414 - val_loss: 1.0985
## Epoch 5/50
## 804/804 - 1s - 1ms/step - accuracy: 0.3396 - loss: 1.0985 - val_accuracy:
```

```
y: 0.3341 - val_loss: 1.0987
## Epoch 6/50
## 804/804 - 1s - 2ms/step - accuracy: 0.3374 - loss: 1.0986 - val_accurac
y: 0.3301 - val_loss: 1.0988
## Epoch 7/50
## 804/804 - 1s - 2ms/step - accuracy: 0.3438 - loss: 1.0983 - val_accurac
y: 0.3383 - val_loss: 1.0987
## Epoch 8/50
## 804/804 - 1s - 2ms/step - accuracy: 0.3444 - loss: 1.0979 - val_accurac
y: 0.3360 - val_loss: 1.0986
## Epoch 9/50
## 804/804 - 1s - 2ms/step - accuracy: 0.3446 - loss: 1.0977 - val_accurac
y: 0.3357 - val_loss: 1.0986
## Epoch 10/50
## 804/804 - 1s - 2ms/step - accuracy: 0.3439 - loss: 1.0979 - val_accurac
y: 0.3385 - val_loss: 1.0985
## Epoch 11/50
## 804/804 - 1s - 2ms/step - accuracy: 0.3495 - loss: 1.0976 - val_accurac
y: 0.3410 - val_loss: 1.0986
## Epoch 12/50
## 804/804 - 1s - 2ms/step - accuracy: 0.3502 - loss: 1.0971 - val_accurac
y: 0.3403 - val_loss: 1.0991
## Epoch 13/50
## 804/804 - 1s - 1ms/step - accuracy: 0.3518 - loss: 1.0972 - val_accurac
y: 0.3371 - val_loss: 1.0987
## Epoch 14/50
## 804/804 - 1s - 2ms/step - accuracy: 0.3538 - loss: 1.0968 - val_accurac
y: 0.3363 - val_loss: 1.0988
## Epoch 15/50
## 804/804 - 1s - 1ms/step - accuracy: 0.3534 - loss: 1.0966 - val_accurac
y: 0.3371 - val_loss: 1.0987
## Epoch 16/50
## 804/804 - 1s - 1ms/step - accuracy: 0.3546 - loss: 1.0967 - val_accurac
y: 0.3382 - val_loss: 1.0989
## Epoch 17/50
## 804/804 - 1s - 1ms/step - accuracy: 0.3573 - loss: 1.0970 - val_accurac
y: 0.3430 - val_loss: 1.0987
## Epoch 18/50
## 804/804 - 1s - 2ms/step - accuracy: 0.3556 - loss: 1.0964 - val_accurac
y: 0.3394 - val_loss: 1.0987
## Epoch 19/50
## 804/804 - 1s - 2ms/step - accuracy: 0.3529 - loss: 1.0963 - val_accurac
y: 0.3413 - val_loss: 1.0988
## Epoch 20/50
## 804/804 - 1s - 1ms/step - accuracy: 0.3527 - loss: 1.0959 - val_accurac
y: 0.3427 - val_loss: 1.0989
## Epoch 21/50
## 804/804 - 1s - 1ms/step - accuracy: 0.3526 - loss: 1.0959 - val_accurac
y: 0.3428 - val_loss: 1.0990
## Epoch 22/50
## 804/804 - 1s - 1ms/step - accuracy: 0.3563 - loss: 1.0951 - val_accurac
y: 0.3397 - val_loss: 1.0995
## Epoch 23/50
## 804/804 - 1s - 1ms/step - accuracy: 0.3596 - loss: 1.0947 - val_accurac
```

```
y: 0.3385 - val_loss: 1.0994
## Epoch 24/50
## 804/804 - 1s - 1ms/step - accuracy: 0.3606 - loss: 1.0952 - val_accurac
y: 0.3341 - val_loss: 1.0994
## Epoch 25/50
## 804/804 - 1s - 1ms/step - accuracy: 0.3560 - loss: 1.0951 - val_accurac
y: 0.3394 - val_loss: 1.1000
## Epoch 26/50
## 804/804 - 1s - 1ms/step - accuracy: 0.3608 - loss: 1.0950 - val_accurac
y: 0.3357 - val_loss: 1.0994
## Epoch 27/50
## 804/804 - 1s - 1ms/step - accuracy: 0.3617 - loss: 1.0952 - val_accurac
y: 0.3378 - val_loss: 1.0995
## Epoch 28/50
## 804/804 - 1s - 1ms/step - accuracy: 0.3605 - loss: 1.0943 - val_accurac
y: 0.3455 - val_loss: 1.0997
## Epoch 29/50
## 804/804 - 1s - 1ms/step - accuracy: 0.3563 - loss: 1.0953 - val_accurac
y: 0.3396 - val_loss: 1.0996
## Epoch 30/50
## 804/804 - 1s - 1ms/step - accuracy: 0.3633 - loss: 1.0949 - val_accurac
y: 0.3361 - val_loss: 1.1000
## Epoch 31/50
## 804/804 - 1s - 1ms/step - accuracy: 0.3624 - loss: 1.0943 - val_accurac
y: 0.3405 - val_loss: 1.1002
## Epoch 32/50
## 804/804 - 1s - 1ms/step - accuracy: 0.3605 - loss: 1.0939 - val_accurac
y: 0.3461 - val_loss: 1.1002
## Epoch 33/50
## 804/804 - 1s - 1ms/step - accuracy: 0.3620 - loss: 1.0944 - val_accurac
y: 0.3378 - val_loss: 1.1004
## Epoch 34/50
## 804/804 - 1s - 1ms/step - accuracy: 0.3597 - loss: 1.0945 - val_accurac
y: 0.3386 - val_loss: 1.1005
## Epoch 35/50
## 804/804 - 1s - 1ms/step - accuracy: 0.3658 - loss: 1.0935 - val_accurac
y: 0.3419 - val_loss: 1.1005
## Epoch 36/50
## 804/804 - 1s - 1ms/step - accuracy: 0.3662 - loss: 1.0935 - val_accurac
y: 0.3427 - val_loss: 1.1011
## Epoch 37/50
## 804/804 - 1s - 1ms/step - accuracy: 0.3584 - loss: 1.0949 - val_accurac
y: 0.3364 - val_loss: 1.1010
## Epoch 38/50
## 804/804 - 1s - 1ms/step - accuracy: 0.3630 - loss: 1.0933 - val_accurac
y: 0.3382 - val_loss: 1.1014
## Epoch 39/50
## 804/804 - 1s - 1ms/step - accuracy: 0.3679 - loss: 1.0938 - val_accurac
y: 0.3344 - val_loss: 1.1001
## Epoch 40/50
## 804/804 - 1s - 1ms/step - accuracy: 0.3654 - loss: 1.0933 - val_accurac
y: 0.3294 - val_loss: 1.1010
## Epoch 41/50
## 804/804 - 1s - 1ms/step - accuracy: 0.3641 - loss: 1.0931 - val_accurac
```

```

y: 0.3282 - val_loss: 1.1010
## Epoch 42/50
## 804/804 - 1s - 1ms/step - accuracy: 0.3618 - loss: 1.0932 - val_accuracy
y: 0.3350 - val_loss: 1.1013
## Epoch 43/50
## 804/804 - 1s - 1ms/step - accuracy: 0.3649 - loss: 1.0939 - val_accuracy
y: 0.3330 - val_loss: 1.1013
## Epoch 44/50
## 804/804 - 1s - 1ms/step - accuracy: 0.3671 - loss: 1.0929 - val_accuracy
y: 0.3396 - val_loss: 1.1015
## Epoch 45/50
## 804/804 - 1s - 2ms/step - accuracy: 0.3602 - loss: 1.0944 - val_accuracy
y: 0.3388 - val_loss: 1.1008
## Epoch 46/50
## 804/804 - 1s - 2ms/step - accuracy: 0.3656 - loss: 1.0933 - val_accuracy
y: 0.3389 - val_loss: 1.1007
## Epoch 47/50
## 804/804 - 1s - 1ms/step - accuracy: 0.3675 - loss: 1.0930 - val_accuracy
y: 0.3343 - val_loss: 1.1010
## Epoch 48/50
## 804/804 - 1s - 1ms/step - accuracy: 0.3661 - loss: 1.0930 - val_accuracy
y: 0.3349 - val_loss: 1.1012
## Epoch 49/50
## 804/804 - 1s - 1ms/step - accuracy: 0.3653 - loss: 1.0934 - val_accuracy
y: 0.3310 - val_loss: 1.1014
## Epoch 50/50
## 804/804 - 1s - 1ms/step - accuracy: 0.3662 - loss: 1.0927 - val_accuracy
y: 0.3333 - val_loss: 1.1011

```

Analysis (Accuracy plot - Top): - Training accuracy hovers around 33% to 35% - Validation accuracy is similar to training accuracy but is slightly lower with both curves being mostly flat and only have a slight upward drift - Hence, model is not overfitting as training and validation accuracy are very close - Shows that the model is not much better than random guessing (3 classes of admission types, so about 33.3% if randomly guess) - So, neural network is not learning meaningful patterns

Analysis (Loss plot - Bottom): - Training loss drops sharply initially then slowly decreases - Validation loss stays almost the same at around 1.1 - Other than the initial gap, most of the gap between the training and validation loss is small - This indicates that initial learning does happen, but the model quickly levels - Though since there is no divergence between the training and validation loss, there unlikely to have any overfitting

## Appendix 6.8: Evaluating Deep Learning Model On Test Data

```

model %>% evaluate(x_test, y_test_cat)

## 251/251 - 0s - 1ms/step - accuracy: 0.3375 - loss: 1.0998

## $accuracy
## [1] 0.3375249
##
## $loss
## [1] 1.099841

```

```

log(3)
## [1] 1.098612

```

Analysis: - Neural network shows an accuracy of about 33.5%, which means it correctly predicts about 33.5% of cases in the test set (close to random guessing) - For the 3 classes, a completely random guess would give loss of 1.099, which is about  $\log(3)$ . This indicates that the model is not learning useful patterns from the data

#### *Appendix 5.6.9: Prediction And Evaluation*

```

# Predict probabilities
pred_prob <- model %>% predict(x_test)

## 251/251 - 0s - 840us/step

# Convert probabilities into class indices starting at 0
pred_class <- apply(pred_prob, 1, which.max) - 1

# Convert indices to factors
pred_class <- factor(
  levels(y_test)[pred_class + 1],
  levels = levels(y_test)
)

# Overall evaluation
confusionMatrix(pred_class, y_test)

## Confusion Matrix and Statistics
##
##             Reference
## Prediction  Elective Emergency Urgent
##   Elective      724       710     767
##   Emergency    324       284     300
##   Urgent        1616      1604    1703
##
## Overall Statistics
##
##                 Accuracy : 0.3375
##                           95% CI : (0.3272, 0.348)
##   No Information Rate : 0.3449
##   P-Value [Acc > NIR] : 0.9189
##
##                 Kappa : -0.002
##
## McNemar's Test P-Value : <2e-16
##
## Statistics by Class:
##
##             Class: Elective Class: Emergency Class: Urgent
## Sensitivity          0.27177      0.10931      0.6148
## Specificity          0.72485      0.88517      0.3881
## Pos Pred Value       0.32894      0.31278      0.3459
## Neg Pred Value       0.66730      0.67518      0.6568

```

## Prevalence	0.33167	0.32346	0.3449
## Detection Rate	0.09014	0.03536	0.2120
## Detection Prevalence	0.27403	0.11305	0.6129
## Balanced Accuracy	0.49831	0.49724	0.5014

Analysis: - Overall accuracy is about 33.5% - Kappa gives 0.0092, means that neural network is not any better than random guessing - Though model does predict emergency more often as it is more sensitive, but precision of guessing is low, which may lead to missclassifications